



Introduction to PL/pgSQL



Procedural Language Overview

- PostgreSQL allows user-defined functions to be written in a variety of procedural languages. The database server has no built-in knowledge about how to interpret the function's source text. Instead, the task is passed to a handler that knows the details of that particular language.
- PostgreSQL currently supports several standard procedural languages
 - PL/pgSQL
 - PL/Tcl
 - PL/Perl
 - PL/Python
 - PL/Java
 - And many more

What is PL/pgSQL

- PL/pgSQL is the procedural extension to SQL with features of programming languages
- Data Manipulation and Query statements of SQL are included within procedural units of code
- Allows using general programming tools with SQL, for example: loops, conditions, functions, etc.
- This allows a lot more freedom than general SQL, and is lighter-weight than calling from a client program

How PL/pgSQL works

- PL/pgSQL is like every other “loadable, procedural language.”
- When a PL function is executed, the fmgr loads the language handler and calls it.
- The language handler then interprets the contents of the pg_proc entry for the function (proargtypes, prorettype, prosrc).

How PL/pgSQL works

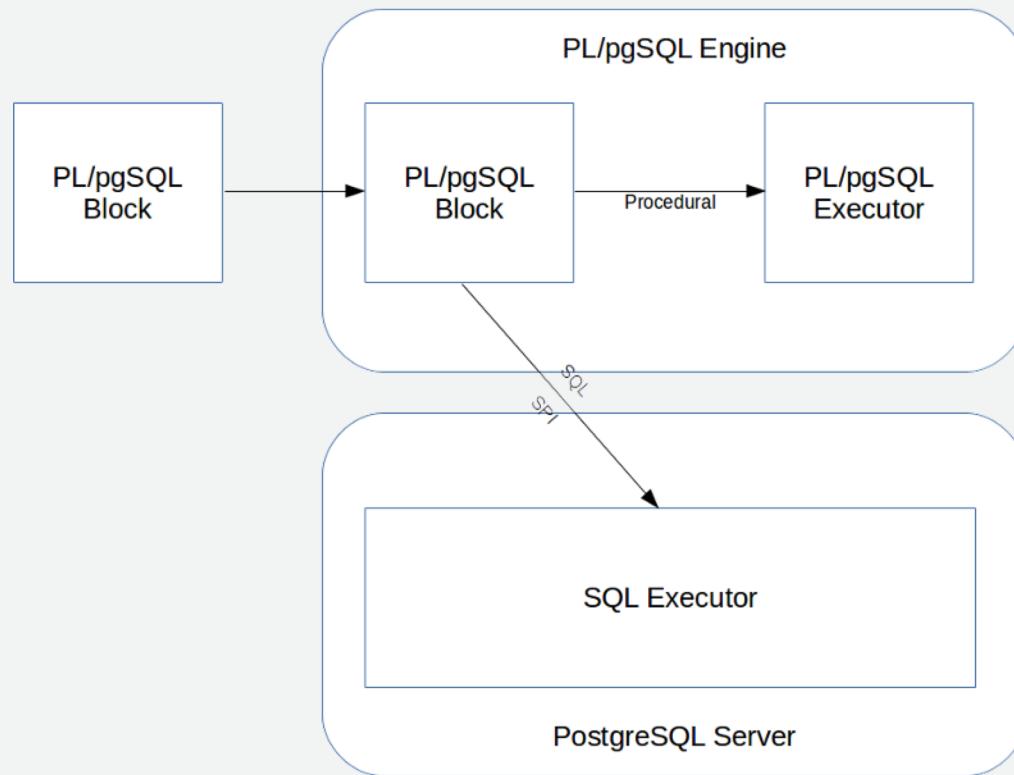
- On the first call of a function in a session, the call handler will “compile” a function statement tree.
- SQL queries in the function are just kept as a string at this point.
- What might look to you like an expression is actually a SELECT query:

```
my_variable := some_parameter * 100;
```

How PL/pgSQL works

- The PL/pgSQL statement tree is very similar to a PostgreSQL execution tree.
- The call handler then executes that statement tree.
- On the first execution of a statement node, that has an SQL query in it, that query is prepared via SPI.
- The prepared plan is then executed for every invocation of that statement in the current session.

PL/pgSQL Environment



Kinds of PL/pgSQL Blocks

The basic unit in any PL/pgSQL code is a BLOCK. All PL/pgSQL code is composed of a single block or blocks that occur either sequentially or nested within another block. There are two kinds of blocks:

- **Anonymous blocks (DO)**
 - Generally constructed dynamically and executed only once by the user. It is sort of a complex SQL statement
- **Named blocks (Functions and Stored Procedures)**
 - Have a name associated with them, are stored in the database, and can be executed repeatably, and can take in parameters

Structure of Anonymous Block

```
DO $$  
[ <<label>> ]  
DECLARE  
/* Declare section (optional). */  
  
BEGIN  
/* Executable section (required). */  
  
EXCEPTION  
/* Exception handling section (optional). */  
  
END [ label ]  
$$;
```

Comments

- There are two types of comments in PL/pgSQL
 - -- starts a comment that extends to the end of the line
 - /* multi-line comments */
- Commenting is necessary to tell people what is intended and why it was done a specific way
- Err on the side of too much commenting

Variables

- Use variables for
 - Temporary storage of data
 - Manipulation of stored values
 - Re-usability
 - Ease of maintenance
- Declared in the declarative section within a block

```
v_last_name      VARCHAR(15);
```

Handling Variables

- Variables declared in the declarations section preceding a block are initialized to their default values every time the block is entered, not only once per function call
- Variables in a declaration section can shadow variables of the same name in an outer block. If the outer block is named with a label, its variables are still available by specifying them as <label>.<varname>

Declarations

Syntax

```
identifier [CONSTANT] datatype [NOT NULL] [:= | = | DEFAULT expr];
```

Examples

```
DECLARE
    v_birthday      DATE;
    v_age           INT NOT NULL = 21;
    v_name          VARCHAR(15) := 'Homer';
    v_magic          CONSTANT NUMERIC := 42;
    v_valid          BOOLEAN DEFAULT TRUE;
```

%TYPE

- Declare variable according to :
 - A database column definition
 - Another previously declared variable

```
identifier table.column_name%TYPE;
```

Example

```
DECLARE
    v_email          users.email%TYPE;
    v_my_email      v_email%TYPE := 'rds-postgres-extensions-
request@amazon.com';
```

%ROWTYPE

- Declare a variable with the type of a ROW of a table

```
identifier table%ROWTYPE;
```

Example

```
DECLARE  
    v_user      users%ROWTYPE;
```

Records

- A record is a type of variable similar to ROWTYPE, but with no predefined structure
- The actual structure of the record is created when the variable is first assigned
- A record is not a true data type, only a place holder

DECLARE

```
r record;
```

Variable Scope

```
DO $$  
DECLARE  
    quantity integer := 30;  
BEGIN  
    RAISE NOTICE 'Quantity here is %', quantity; -- 30  
    quantity := 50;  
    -- Create a subblock  
    DECLARE  
        quantity integer := 80;  
    BEGIN  
        RAISE NOTICE 'Quantity here is %', quantity; -- 80  
    END;  
    RAISE NOTICE 'Quantity here is %', quantity; -- 50  
END  
$$;
```

Qualify an Identifier

```
DO $$  
  << mainblock >>  
DECLARE  
    quantity integer := 30;  
BEGIN  
    RAISE NOTICE 'Quantity here is %', quantity; --30  
    quantity := 50;  
    -- Create a subblock  
    DECLARE  
        quantity integer := 80;  
    BEGIN  
        RAISE NOTICE 'Quantity here is %', mainblock.quantity; --50  
        RAISE NOTICE 'Quantity here is %', quantity; --80  
    END;  
    RAISE NOTICE 'Quantity here is %', quantity; --50  
END  
$$;
```

RAISE

- Reports messages
 - Can be seen by the client if the appropriate level is used

```
RAISE NOTICE 'Calling cs_create_job(%)', v_job_id;
```

Assigning Values

- Use the assignment operator (:= or =)

```
DECLARE
    v_last_name      VARCHAR := 'Smith';
    v_date           DATE;
BEGIN
    v_last_name := lower(v_last_name);
    v_date := to_date('2000-01-01', 'YYYY-MM-DD');
```

SELECT in PL/pgSQL

- Retrieve data from the database with a SELECT statement
- Queries must return only one row
- INTO clause is required

```
DECLARE
    v_first_name    users.first_name%TYPE;
    v_last_name     users.last_name%TYPE;
BEGIN
    SELECT first_name, last_name
        INTO v_first_name, v_last_name
        FROM users
        WHERE user_id = 1;
END
```

INSERT / UPDATE / DELETE

```
DECLARE
    v_forum_name    forums.name%TYPE := 'Hackers';
BEGIN
    INSERT INTO forums (name)
    VALUES (v_forum_name);

    UPDATE forums
        SET moderated = true
        WHERE name = v_forum_name;
END
```

PERFORM

- Evaluate an expression or query but discard the result
- Frequently used when executing maintenance commands

BEGIN

```
PERFORM create_partition('moderation_log', '2016-06');
```

END

Structure of Named Blocks

```
CREATE FUNCTION [ function_name ] ()
RETURNS [return_type] $$  
[ <<label>> ]
DECLARE
/* Declare section (optional). */

BEGIN
/* Executable section (required). */

EXCEPTION
/* Exception handling section (optional). */

END [ label ]
$$ LANGUAGE plpgsql;
```

Function Example

```
CREATE FUNCTION get_user_count()
    RETURNS integer
AS $$

DECLARE
    v_count      integer;
BEGIN
    SELECT count(*)
        INTO v_count
    FROM users;

    RETURN v_count;
END
$$ LANGUAGE plpgsql;
```

Dollar Quoting

- The tag \$\$ denotes the start and end of a string
- Optionally can have a non-empty tag as part of the quote
 - \$_\$
 - \$abc\$
- Can be used to prevent unnecessary escape characters throughout the string

```
$function$  
BEGIN  
    RETURN ($1 ~ $q$[\t\r\n\v\\]$q$);  
END;  
$function$
```

Function Parameters

- One or more parameters can be used
- Parameter names are optional, but highly recommended

```
CREATE FUNCTION get_user_name(varchar, p_last_name varchar)
    RETURNS varchar AS $$  
DECLARE
    v_first_name      varchar;
    v_name            varchar;
BEGIN
    v_first_name := $1;
    SELECT name INTO v_name FROM users
        WHERE first_name = v_first_name AND last_name = p_last_name
        LIMIT 1;

    RETURN v_name;
END
$$ LANGUAGE plpgsql;
```

Default Parameters

- Parameters can have a default value
- This essentially makes them optional parameters

```
CREATE FUNCTION get_user_count(p_active boolean DEFAULT true)
RETURNS integer AS $$  
DECLARE
    v_count      integer;
BEGIN
    SELECT count(*) INTO v_count
    FROM users
    WHERE active = p_active;

    RETURN v_count;
END
$$ LANGUAGE plpgsql;
```

Assertions

- A convenient shorthand for inserting debugging checks
- Can be controlled by plpgsql.check_asserts variable

```
CREATE FUNCTION get_user_count(p_active boolean DEFAULT true)
    RETURNS integer AS $$

DECLARE
    v_count      integer;
BEGIN
    ASSERT p_active IS NOT NULL;

    SELECT count(*) INTO v_count
        FROM users
        WHERE active = p_active;

    RETURN v_count;
END
$$ LANGUAGE plpgsql;
```

PL/pgSQL Control Structures



Control the Flow

- The logical flow of statements can be changed using conditional IF statements and loop control structures
 - Conditional Structures
 - Loop Structures

IF Statements

IF-THEN

```
IF boolean-expression THEN  
    statements  
END IF;
```

IF-THEN-ELSE

```
IF boolean-expression THEN  
    statements  
ELSE  
    statements  
END IF;
```

Nested IF Statements

```
IF boolean-expression THEN
    IF boolean-expression THEN
        statements
    END IF;
ELSE
    statements
END IF;
```

ELSIF Statements

- A sequence of statements based on multiple conditions

```
IF number = 0 THEN
    result := 'zero';
ELSIF number > 0 THEN
    result := 'positive';
ELSIF number < 0 THEN
    result := 'negative';
ELSE
    -- the only other possibility is that number is null
    result := 'NULL';
END IF;
```

CASE Statements

- Used for complex conditionals
- Allows a variable to be tested for equality against a list of values

```
BEGIN
    CASE status
        WHEN 'Pending' THEN RAISE NOTICE 'PENDING';
        WHEN 'Accepted' THEN RAISE NOTICE 'ACCEPTED';
        WHEN 'Declined' THEN RAISE NOTICE 'DECLINED';
        WHEN 'Blocked' THEN RAISE NOTICE 'BLOCKED';
        ELSE RAISE NOTICE 'UNKNOWN';
    END CASE;
END
```

Searched CASE Statements

- Each WHEN clause sequentially evaluated until a TRUE is evaluated
- Subsequent WHEN expressions are not evaluated

```
BEGIN
CASE
    WHEN x BETWEEN 0 AND 10 THEN
        RAISE NOTICE 'Value is between zero and ten';
    WHEN x BETWEEN 11 AND 20 THEN
        RAISE NOTICE 'Value is between eleven and twenty';
END CASE;
$$;
```

FOUND

- FOUND, which is of type boolean, starts out false within each PL/pgSQL function call
- It is set by each of the following types of statements:
 - A SELECT INTO statement sets FOUND true if it returns a row, false if no row is returned
 - A PERFORM statement sets FOUND true if it produces (and discards) a row, false if no row is produced
 - UPDATE, INSERT, and DELETE statements set FOUND true if at least one row is affected, false if no row is affected
 - A FETCH statement sets FOUND true if it returns a row, false if no row is returned.
 - A FOR statement sets FOUND true if it iterates one or more times, else false.

FOUND

```
DECLARE
    v_first_name    users.first_name%TYPE;
    v_last_name     users.last_name%TYPE;
BEGIN
    SELECT first_name, last_name
        INTO v_first_name, v_last_name
        FROM users
       WHERE user_id = 1;

    IF FOUND THEN
        RAISE NOTICE 'User Found';
    ELSE
        RAISE NOTICE 'User Not Found';
    END IF;
END
```

Loop Structures

- Unconstrained Loop
- WHILE Loop
- FOR Loop
- FOREACH Loop

Unconstrained Loops

- Allows execution of its statements at least once, even if the condition already met upon entering the loop

```
LOOP
    -- some computations
    IF count > 0 THEN
        EXIT;  -- exit loop
    END IF;
END LOOP;
LOOP
    -- some computations
    EXIT WHEN count > 0;  -- same result as previous example
END LOOP;
```

CONTINUE

```
CONTINUE [ label ] [ WHEN expression ];
```

- If no label is given, the next iteration of the innermost loop is begun
- If WHEN is specified, the next iteration of the loop is begun only if expression is true.
Otherwise, control passes to the statement after CONTINUE
- CONTINUE can be used with all types of loops; it is not limited to use with unconstrained loops.

```
LOOP
```

```
-- some computations  
EXIT WHEN count > 100;  
CONTINUE WHEN count < 50;  
-- some computations for count IN [50 .. 100]
```

```
END LOOP;
```

WHILE Loops

```
WHILE condition LOOP  
    statement1..;  
END LOOP;
```

- Repeats a sequence of statements until the controlling condition is no longer TRUE
- Condition is evaluated at the beginning of each iteration

```
WHILE NOT done LOOP  
    -- some computations here  
END LOOP;
```

FOR Loops

```
FOR <loop_counter> IN [REVERSE] <low bound>..<high bound> LOOP  
    -- some computations here  
END LOOP;
```

- Use a FOR loop to shortcut the test for the number of iterations.
- Do not declare the counter; it is declared implicitly

```
DO $$  
BEGIN  
    FOR i IN 1..10 LOOP  
        RAISE NOTICE 'value: %', i;  
    END LOOP;  
END  
$$;
```

Looping Over Results

- For loops can directly use a query result

```
DECLARE
    r record;
BEGIN
    FOR r IN SELECT email FROM users LOOP
        RAISE NOTICE 'Email: %', r.email;
    END LOOP;
END
```

Looping Over Results

- The last row is still accessible after exiting the loop

```
DECLARE
    r record;
BEGIN
    FOR r IN SELECT email FROM users LOOP
        RAISE NOTICE 'Email: %', r.email;
    END LOOP;
    RAISE NOTICE 'Email: %', r.email;
END
```

Looping Over Results

- Looping over dynamic SQL
- Re-planned each time it is executed

```
DECLARE
    rec RECORD;
    sql text;
BEGIN
    sql := 'SELECT email FROM users';
    FOR rec IN EXECUTE sql LOOP
        RAISE NOTICE 'Email: %', rec.email;
    END LOOP;
END
```

Looping Over Arrays

- Uses the FOREACH statement

```
DECLARE
    users      varchar[ ] := ARRAY[ 'Mickey', 'Donald', 'Minnie' ];
    v_user     varchar;
BEGIN
    FOREACH v_user IN ARRAY users LOOP
        RAISE NOTICE 'User: %', v_user;
    END LOOP;
END
```

Looping Over Arrays

- Use the SLICE syntax to iterate over multiple dimensions

```
DECLARE
    users      varchar[ ];
    v_dim      varchar[ ];
BEGIN
    users := ARRAY[ARRAY[ 'Mickey' , 'Donald' ], ARRAY[ 'Mouse' ,
'Duck' ]];
    FOREACH v_dim SLICE 1 IN ARRAY users LOOP
        RAISE NOTICE 'Dimension: %', v_dim;
    END LOOP;
END
```

Nested Loops

- Nest loops to multiple levels
 - Use labels to distinguish between blocks
 - Exit the outer loop with the EXIT statement that references the label

```
BEGIN
    <<Outer_loop>>
    LOOP
        v_counter := v_counter + 1;
    EXIT WHEN v_counter > 10; -- leaves both loops
    <<Inner_loop>>
    LOOP
        EXIT Outer_loop WHEN total_done = 'YES';
        -- leaves both loops
        EXIT WHEN inner_done = 'YES';
        -- leaves inner loop only
        END LOOP Inner_loop;
    END LOOP Outer_loop;
END
```

Dynamic SQL



Dynamic SQL

- A programming methodology for generating and running SQL statements at run time

- Useful for:

- Ad-hoc query systems
- DDL and database maintenance

```
EXECUTE command-string [ INTO target ] [ USING expression [, ...] ];
```

Dynamic SQL - CAUTION

- There is no plan caching for commands executed via EXECUTE
 - The command is planned each time it is run
- Open to SQL injection attacks
 - All incoming parameters need to be validated
 - Bind the parameters to the command instead of generating the string

Execute

```
CREATE FUNCTION grant_select(p_table varchar, p_role varchar)
    RETURNS void AS
$$
DECLARE
    sql      varchar;
BEGIN
    sql := 'GRANT SELECT ON TABLE ' || p_table || ' TO ' || p_role;
    EXECUTE sql;
END
$$ LANGUAGE plpgsql;
```

Note: Do not do this. Validate the parameters first.

Execute Into

```
CREATE FUNCTION get_connection_count(p_role varchar)
    RETURNS integer
AS $$

DECLARE
    v_count      integer;
    sql          varchar;

BEGIN
    sql := 'SELECT count(*) FROM pg_stat_activity
            WHERE username = ''' || p_role || '''';
    EXECUTE sql INTO v_count;

    RETURN v_count;
END
$$ LANGUAGE plpgsql;
```

Note: Do not do this. Validate the parameters first.

Execute Using

```
CREATE FUNCTION get_connection_count(p_role varchar)
    RETURNS integer
AS $$

DECLARE
    v_count      integer;
    sql          varchar;

BEGIN
    sql := 'SELECT count(*) FROM pg_stat_activity
            WHERE username = $1';
    EXECUTE sql INTO v_count USING p_role;

    RETURN v_count;
END
$$ LANGUAGE plpgsql;
```

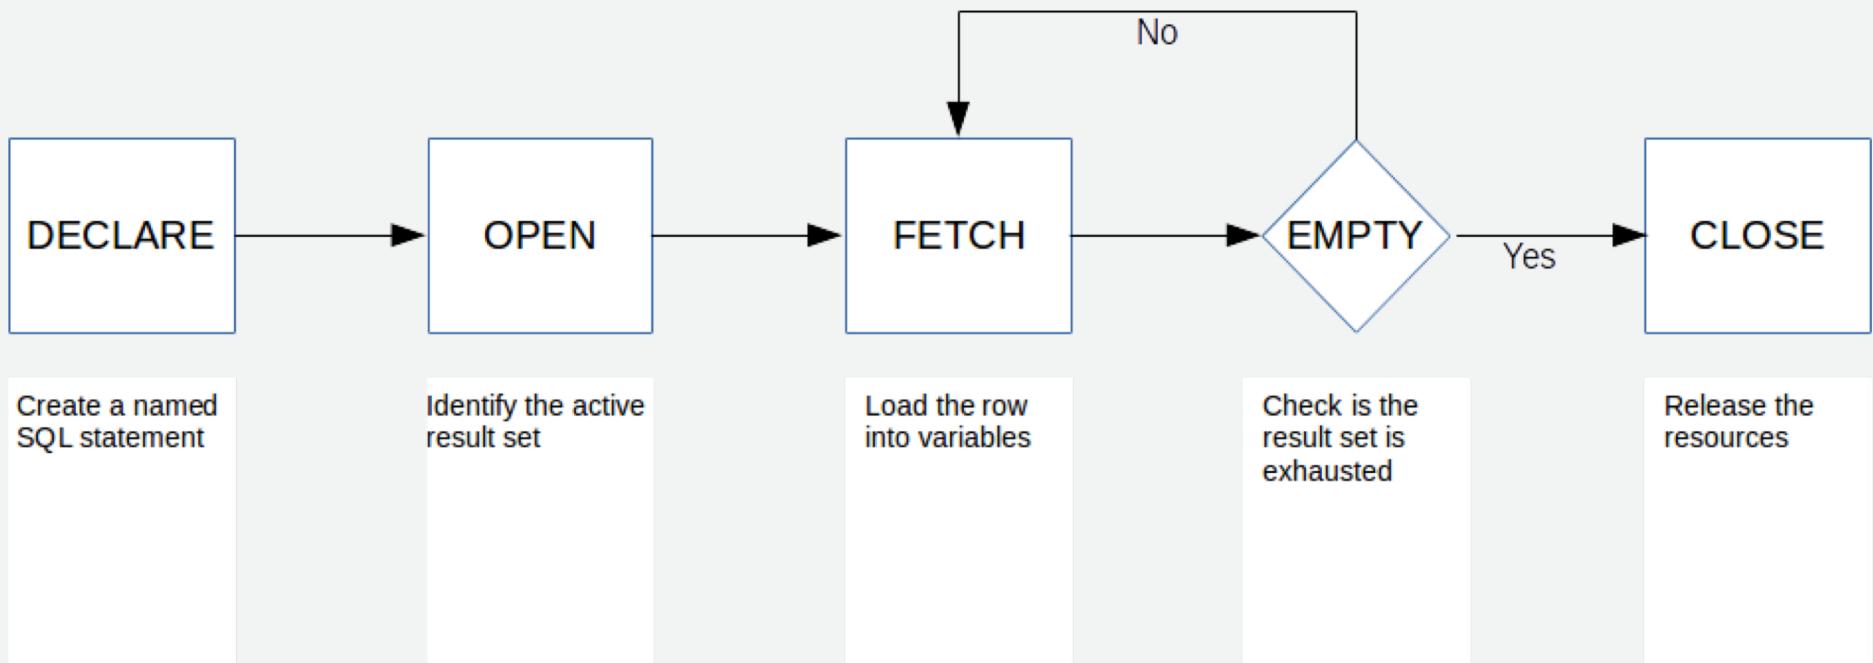
PL/pgSQL Cursors



Cursors

- Every SQL statement executed by PostgreSQL has an individual cursor associated with it
 - Implicit cursors: Declared for all DML and PL/pgSQL SELECT statements
 - Explicit cursors: Declared and named by the programmer
- Use CURSOR to individually process each row returned by a multiple-row SELECT Statement

Cursor Flow



Declaring Cursors

- A cursor must be declared as a variable
 - Use the SCROLL keyword to move backwards through a cursor

```
name [ [ NO ] SCROLL ] CURSOR [ ( arguments ) ] FOR query;  
DECLARE  
    curs1 refcursor;  
    curs2 CURSOR FOR SELECT * FROM tenk1;  
    curs3 CURSOR (key integer) FOR SELECT *  
                      FROM tenk1  
                      WHERE unique1 = key;
```

Opening Cursors

- The OPEN method to use is dependant on the way it was declared

```
OPEN curs1 FOR SELECT * FROM foo WHERE key = mykey;  
OPEN cur2;
```

```
OPEN curs3(42);  
OPEN curs3 (key := 42);
```

Fetching Data

- **FETCH returns the next row**

```
FETCH curs2 INTO foo, bar, baz;
```

- **FETCH can also move around the cursor**

```
FETCH LAST FROM curs3 INTO x, y;
```

Fetching Data

```
CREATE FUNCTION grant_select(p_role varchar) RETURNS void AS $$  
DECLARE  
    sql      varchar;  
    r        record;  
    tbl_cursor CURSOR FOR SELECT schemaname, relname  
                         FROM pg_stat_user_tables;  
BEGIN  
    OPEN tbl_cursor;  
    LOOP  
        FETCH tbl_cursor INTO r;  
        EXIT WHEN NOT FOUND;  
        sql := 'GRANT SELECT ON TABLE ' || r.schemaname ||  
               '.' || r.relname || ' TO ' || p_role;  
        EXECUTE sql;  
    END LOOP;  
    CLOSE tbl_cursor;  
END $$ LANGUAGE plpgsql;
```

PL/pgSQL Returning Data



Returning Scalars

```
CREATE FUNCTION
get_connection_count()
    RETURNS integer AS $$

DECLARE
    v_count      integer;
BEGIN
    SELECT count(*) INTO
v_count
        FROM pg_stat_activity;

    RETURN v_count;
END
$$ LANGUAGE plpgsql;
```

```
SELECT
get_connection_count();
get_connection_count
-----
(1 row)
```

11

Returning Nothing

- Some functions do not need a return value
 - This is usually a maintenance function of some sort such as creating partitions or data purging
 - Starting in PostgreSQL 11, Stored Procedures can be used in these cases
- Return VOID

```
CREATE FUNCTION purge_log()
RETURNS void AS
$$
BEGIN
    DELETE FROM moderation_log
    WHERE log_date < now() - '90 days'::interval;
END
$$ LANGUAGE plpgsql;
```

Returning Sets

- Functions can return a result set
- Use SETOF
- Use RETURN NEXT
 - RETURN NEXT does not actually return from the function
 - Successive RETURN NEXT commands build a result set
- A final RETURN exits the function

Returning Sets

```
CREATE FUNCTION fibonacci(num integer) RETURNS SETOF integer AS $$  
DECLARE  
    a int := 0;  
    b int := 1;  
BEGIN  
    IF (num <= 0)  
        THEN RETURN;  
    END IF;  
  
    RETURN NEXT a;  
    LOOP  
        EXIT WHEN num <= 1;  
        RETURN NEXT b;  
        num = num - 1;  
        SELECT b, a + b INTO a, b;  
    END LOOP;  
END; $$ language plpgsql;
```

Returning Records

```
CREATE FUNCTION get_oldest_session()
    RETURNS record AS
$$
DECLARE
    r      record;
BEGIN
    SELECT *
        INTO r
        FROM pg_stat_activity
        WHERE usename = SESSION_USER
        ORDER BY backend_start DESC
        LIMIT 1;

    RETURN r;
END
$$ LANGUAGE plpgsql;
```

Returning Records

- Using a generic record type requires the structure to be defined at run time

```
# SELECT * FROM get_oldest_session();
ERROR: a column definition list is required for functions ...
LINE 1: SELECT * FROM get_oldest_session();
```

```
SELECT * FROM get_oldest_session()
AS (a oid, b name, c integer, d oid, e name, f text, g inet,
    h text, i integer, j timestamptz, k timestamptz,
    l timestamptz, m timestamptz, n boolean, o text, p xid,
    q xid, r text);
```

Returning Records

- All tables and views automatically have corresponding type definitions so they can be used as return types

```
CREATE FUNCTION get_oldest_session() RETURNS pg_stat_activity AS $$  
DECLARE  
    r      record;  
BEGIN  
    SELECT *  
    INTO r  
    FROM pg_stat_activity  
    WHERE username = SESSION_USER  
    ORDER BY backend_start DESC  
    LIMIT 1;  
  
    RETURN r;  
END $$ LANGUAGE plpgsql;
```

Returning Sets of Records

- Many times, a subset of the table data is needed
- A view can be used to define the necessary structure

```
CREATE VIEW running_queries AS
SELECT CURRENT_TIMESTAMP - query_start AS runtime, pid,
       username, waiting, query
  FROM pg_stat_activity
 ORDER BY 1 DESC
LIMIT 10;
```

Returning Sets of Records

- RETURN QUERY can be used to simplify the function

```
CREATE FUNCTION running_queries(p_rows int, p_len int DEFAULT 50)
    RETURNS SETOF running_queries AS
$$
BEGIN
    RETURN QUERY SELECT runtime, pid, username, waiting,
        substring(query,1,p_len) as query
        FROM running_queries
        ORDER BY 1 DESC
        LIMIT p_rows;
END
$$ LANGUAGE plpgsql;
```

OUT Parameters

- Used to return structured information
- RETURNS is optional, but must be record if included

```
CREATE FUNCTION active_locks(OUT p_excl int, OUT p_share int)
```

OUT Parameters

```
CREATE FUNCTION active_locks(OUT p_excl int, OUT p_share int) AS $$  
DECLARE  
    r      record;  
BEGIN  
    p_excl := 0;  
    p_share := 0;  
    FOR r IN SELECT l.mode  
        FROM pg_locks l, pg_stat_activity a  
        WHERE a.pid = l.pid AND a.username = SESSION_USER  
    LOOP  
        IF r.mode = 'ExclusiveLock' THEN  
            p_excl := p_excl + 1;  
        ELSIF r.mode = 'ShareLock' THEN  
            p_share := p_share + 1;  
        END IF;  
    END LOOP;  
END $$ LANGUAGE plpgsql;
```

OUT Parameters

- TIP: Think in sets not loops when writing functions for better performance
- NOTE: Use “OR REPLACE” when updating functions

```
CREATE OR REPLACE FUNCTION active_locks(OUT p_excl int,
                                         OUT p_share int)
AS $$  
BEGIN  
    SELECT sum(CASE l.mode WHEN 'ExclusiveLock' THEN 1 ELSE 0 END),
           sum(CASE l.mode WHEN 'ShareLock' THEN 1 ELSE 0 END)
      INTO p_excl, p_share
     FROM pg_locks l, pg_stat_activity a
    WHERE a.pid = l.pid
      AND a.username = SESSION_USER;  
  
END  
$$ LANGUAGE plpgsql;
```

Structured Record Sets

- Use OUT parameters and SETOF record

```
CREATE FUNCTION all_active_locks(OUT p_lock_mode varchar,
                                OUT p_count int)
RETURNS SETOF record AS $$  
DECLARE  
    r      record;  
BEGIN  
    FOR r IN SELECT l.mode, count(*) as k  
        FROM pg_locks l, pg_stat_activity a  
        WHERE a.pid = l.pid  
        AND a.username = SESSION_USER  
        GROUP BY 1  
    LOOP  
        p_lock_mode := r.mode;  
        p_count := r.k;  
        RETURN NEXT;  
    ...  
END$$;
```

Structured Record Sets

- Can return a TABLE

```
CREATE FUNCTION all_active_locks()
    RETURNS TABLE (p_lock_mode varchar, p_count int) AS $$  
DECLARE
    r      record;
BEGIN
    FOR r IN SELECT l.mode, count(*) as k
        FROM pg_locks l, pg_stat_activity a
        WHERE a.pid = l.pid
        AND a.username = SESSION_USER
        GROUP BY 1
LOOP
    p_lock_mode := r.mode;
    p_count := r.k;
    RETURN NEXT;
...

```

Refcursors

- A cursor can be returned for large result sets
- The only way to return multiple result sets from a function

```
CREATE FUNCTION active_info(OUT p_queries refcursor,  
                           OUT p_locks refcursor)
```

Refcursors

```
CREATE FUNCTION active_info(OUT p_queries refcursor,
                            OUT p_locks refcursor)
AS $$

BEGIN
    OPEN p_queries FOR SELECT runtime, pid, username, waiting,
                           substring(query,1,50) as query
                           FROM running_queries
                           ORDER BY 1 DESC;

    OPEN p_locks FOR SELECT l.mode, count(*) as k
                   FROM pg_locks l, pg_stat_activity a
                   WHERE a.pid = l.pid
                     AND a.username = SESSION_USER
                   GROUP BY 1;

END
$$ LANGUAGE plpgsql;
```

Handling Meta Information and Exceptions

Meta Information

- Information about the last command run inside of a function
- Several available values
 - ROW_COUNT
 - RESULT_OID
 - PG_CONTEXT

```
GET DIAGNOSTICS variable { = | := } item [ , ... ];
```

Meta Information

```
CREATE OR REPLACE FUNCTION purge_log()
    RETURNS void AS
$$
DECLARE
    l_rows      int;
BEGIN
    DELETE FROM moderation_log
        WHERE log_date < now() - '90 days'::interval;

    GET DIAGNOSTICS l_rows = ROW_COUNT;
    RAISE NOTICE 'Deleted % rows from the log', l_rows;
END
$$ LANGUAGE plpgsql;
```

Exceptions

- An exception is an identifier in PL/pgSQL that is raised during execution
- It is raised when an error occurs or explicitly by the function
- It is either handled in the EXCEPTION block or propagated to the calling environment

```
[ DECLARE ]  
BEGIN  
    Exception/Error is Raised  
EXCEPTION  
    Error is Trapped  
END
```

Exceptions

- Use the WHEN block inside of the EXCEPTION block to catch specific cases
- Can use the error name or error code in the EXCEPTION block

```
WHEN division_by_zero THEN ...  
WHEN SQLSTATE '22012' THEN ...
```

- Use the special conditions OTHERS as a catch all

```
WHEN OTHERS THEN ...
```

Sample Error Codes

Code	Name
22000	data_exception
22012	division_by_zero
2200B	escape_character_conflict
22007	invalid_datetime_format
22023	invalid_parameter_value
2200M	invalid_xml_document
2200S	invalid_xml_comment
23P01	exclusionViolation

Exceptions

```
CREATE OR REPLACE FUNCTION get_connection_count()
    RETURNS integer AS $$

DECLARE
    v_count      integer;
BEGIN
    SELECT count(*) INTO STRICT v_count
        FROM pg_stat_activity;
    RETURN v_count;
EXCEPTION
    WHEN TOO_MANY_ROWS THEN
        RAISE NOTICE 'More than 1 row returned';
        RETURN -1;
    WHEN OTHERS THEN
        RAISE NOTICE 'Unknown Error';
        RETURN -1;
END $$ LANGUAGE plpgsql;
```

Exception Information

- SQLSTATE Returns the numeric value for the error code.
- SQLERRM Returns the message associated with the error number.

```
DECLARE
    v_count      integer;
    err_num      integer;
    err_msg      varchar;

BEGIN
    . . .
EXCEPTION
    WHEN OTHERS THEN
        err_num  := SQLSTATE;
        err_msg   := SUBSTR(SQLERRM,1,100);
        RAISE NOTICE 'Trapped Error: %', err_msg;
        RETURN -1;
END
```

Exception Information

- The details of an error are usually required when handling
- Use GET STACKED DIAGNOSTICS to return the details

```
GET STACKED DIAGNOSTICS variable { = | := } item [ , ... ];
```

Exception Information

Diagnostic Item
RETURNED_SQLSTATE
COLUMN_NAME
CONSTRAINT_NAME
PG_DATATYPE_NAME
MESSAGE_TEXT
TABLE_NAME
SCHEMA_NAME
PG_EXCEPTION_DETAIL
PG_EXCEPTION_HINT
PG_EXCEPTION_CONTEXT

Propagating Exceptions

- Exceptions can be raised explicitly by the function

```
CREATE OR REPLACE FUNCTION grant_select(p_role varchar)
    RETURNS void AS
$$
DECLARE
    sql      varchar;
    r        record;
    tbl_cursor CURSOR FOR SELECT schemaname, relname
                         FROM pg_stat_user_tables;
BEGIN
    IF NOT EXISTS (SELECT 1 FROM pg_roles
                   WHERE rolname = p_role) THEN
        RAISE EXCEPTION 'Invalid Role: %', p_role;
    END IF;
    ...

```

Exceptions

```
CREATE FUNCTION t1()
    RETURNS void AS $$

DECLARE
    i integer;
BEGIN
    i := 1;
END
$$ LANGUAGE plpgsql;
```

Avg Time: 0.0017ms

```
CREATE FUNCTION t2()
    RETURNS void AS $$

DECLARE
    i integer;
BEGIN
    i := 1;
EXCEPTION
    WHEN OTHERS THEN
        RETURN;
END
$$ LANGUAGE plpgsql;
```

Avg Time: 0.0032ms

PL/pgSQL Triggers



Triggers

- Code that gets executed when an event happens in the database
 - INSERT, UPDATE, DELETE
- Event Triggers fire on DDL
 - CREATE, DROP, ALTER

Use Cases

- Table Partitioning before PostgreSQL 10
- Automatically generate derived column values
- Enforce complex constraints
- Enforce referential integrity across nodes in a distributed database
- Provide transparent event logging
- Provide auditing
- Invalidate cache entries

Structure

- Unlike other databases, a trigger is broken into two pieces
 - Trigger
 - Trigger Function

```
CREATE TRIGGER name
  { BEFORE | AFTER | INSTEAD OF }
  { event [ OR ... ] }
  ON table_name
  [ FOR [ EACH ] { ROW | STATEMENT } ]
  [ WHEN ( condition ) ]
  EXECUTE PROCEDURE function_name ( arguments )
```

Trigger Function

- A function with no parameters that returns TRIGGER

```
CREATE FUNCTION trg() RETURNS trigger AS $$  
BEGIN  
    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;
```

Trigger Events

- Insert
- Update
- Delete
- Truncate

Timing

- **Before**
 - The trigger is fired before the change is made to the table
 - Trigger can modify NEW values
 - Trigger can suppress the change altogether
- **After**
 - The trigger is fired after the change is made to the table
 - Trigger sees final result of row

Frequency

- **For Each Row**
 - The trigger is fired once each time a row is affected
- **For Each Statement**
 - The trigger is fired once each time a statement is executed

Row Triggers

```
CREATE FUNCTION debug_cities_insert()
    RETURNS trigger AS
$$
BEGIN
    RAISE NOTICE 'New city inserted';
    RETURN new;
END
$$
LANGUAGE plpgsql;

CREATE TRIGGER debug_cities_insert_trigger
AFTER INSERT ON cities
FOR EACH ROW EXECUTE PROCEDURE debug_cities_insert();
```

Statement Triggers

```
CREATE FUNCTION stop_truncate()
    RETURNS trigger AS
$$
BEGIN
    RAISE EXCEPTION 'This table can not be truncated';
    RETURN null;
END
$$
LANGUAGE plpgsql;

CREATE TRIGGER stop_truncate_trigger
BEFORE TRUNCATE ON cities
FOR EACH STATEMENT EXECUTE PROCEDURE stop_truncate();
```

Arguments

- NEW
 - Variable holding the new row for INSERT/UPDATE operations in row-level triggers
- OLD
 - Variable holding the old row for UPDATE/DELETE operations in row-level triggers

NEW vs OLD

```
CREATE TABLE audit (
    event_time timestamp NOT NULL,
    user_name  varchar  NOT NULL,
    old_row    json,
    new_row    json
);
```

NEW vs OLD

```
CREATE OR REPLACE FUNCTION audit_trigger()
RETURNS TRIGGER AS $$

BEGIN
    INSERT INTO audit
        VALUES (CURRENT_TIMESTAMP,
                CURRENT_USER,
                row_to_json(OLD),
                row_to_json(NEW));

    RETURN NEW;
END;
$$
LANGUAGE plpgsql;
```

Arguments

- **TG_OP**
 - A string of INSERT, UPDATE, DELETE, or TRUNCATE telling for which operation the trigger was fired
- **TG_NAME**
 - Variable that contains the name of the trigger actually fired
- **TG_WHEN**
 - A string of BEFORE, AFTER, or INSTEAD OF, depending on the trigger's definition
- **TG_LEVEL**
 - A string of either ROW or STATEMENT depending on the trigger's definition

TG_OP

```
CREATE TABLE audit (
    event_time timestamp NOT NULL,
    user_name  varchar  NOT NULL,
    operation   varchar  NOT NULL,
    old_row     json,
    new_row     json
);
```

TG_OP

```
CREATE FUNCTION audit_trigger() RETURNS TRIGGER AS $$  
BEGIN  
    IF (TG_OP = 'DELETE') THEN  
        INSERT INTO audit VALUES  
        (CURRENT_TIMESTAMP, CURRENT_USER,TG_OP, row_to_json(OLD), null);  
        RETURN OLD;  
    ELSIF (TG_OP = 'UPDATE') THEN  
        INSERT INTO audit VALUES  
        (CURRENT_TIMESTAMP, CURRENT_USER,TG_OP,  
         row_to_json(OLD), row_to_json(NEW));  
        RETURN NEW;  
    ELSIF (TG_OP = 'INSERT') THEN  
        INSERT INTO audit VALUES  
        (CURRENT_TIMESTAMP, CURRENT_USER,TG_OP, null, row_to_json(NEW));  
        RETURN NEW;  
    END IF;  
    RETURN NULL;  
END $$ LANGUAGE plpgsql;
```

Arguments

- **TG_TABLE_NAME**
 - The name of the table that caused the trigger invocation.
- **TG_RELNAME**
 - The name of the table that caused the trigger invocation
- **TG_RELID**
 - The object ID of the table that caused the trigger invocation
- **TG_TABLE_SCHEMA**
 - The name of the schema of the table that caused the trigger invocation

TG_TABLE_NAME

```
CREATE TABLE audit (
    event_time timestamp NOT NULL,
    user_name  varchar  NOT NULL,
    operation  varchar  NOT NULL,
    table_name varchar  NOT NULL,
    old_row    json,
    new_row    json
);
```

TG_TABLE_NAME

```
CREATE OR REPLACE FUNCTION audit_trigger() RETURNS TRIGGER AS $$  
BEGIN  
    IF (TG_OP = 'DELETE') THEN  
        INSERT INTO audit  
            VALUES (CURRENT_TIMESTAMP, CURRENT_USER,TG_OP,  
                    TG_TABLE_NAME, row_to_json(OLD), null);  
        RETURN OLD;  
    ELSIF (TG_OP = 'UPDATE') THEN  
        INSERT INTO audit  
            VALUES (CURRENT_TIMESTAMP, CURRENT_USER,TG_OP,  
                    TG_TABLE_NAME, row_to_json(OLD),  
                    row_to_json(NEW));  
        RETURN NEW;  
    ...  
END$$;
```

Arguments

- **TG_NARGS**
 - The number of arguments given to the trigger procedure in the CREATE TRIGGER statement
- **TG_argv[]**
 - The arguments from the CREATE TRIGGER statement

Trigger Use Cases

- **Table Partitioning**
 - Splitting what is logically one large table into smaller physical pieces
- **Used to:**
 - Increase performance
 - Archive data
 - Storage tiering

Table Partitioning before PostgreSQL 10

- Create child tables for each partition

```
CREATE TABLE audit_2014 (
    CHECK ( event_time >= DATE '2014-01-01'
            AND event_time < DATE '2015-01-01' )
) INHERITS (audit);
```

```
CREATE TABLE audit_2015 (
    CHECK ( event_time >= DATE '2015-01-01'
            AND event_time < DATE '2016-01-01' )
) INHERITS (audit);
```

Table Partitioning before PostgreSQL 10

- The trigger function will move the row to the correct child table

```
CREATE OR REPLACE FUNCTION partition_audit_trigger()
RETURNS TRIGGER AS $$

BEGIN
    EXECUTE 'INSERT INTO audit_'
        || to_char(NEW.event_time, 'YYYY')
        || ' VALUES ($1, $2, $3, $4, $5, $6)'
    USING NEW.event_time, NEW.user_name, NEW.operation,
          NEW.table_name, NEW.old_row, NEW.new_row;

    RETURN NULL;
END;
$$
LANGUAGE plpgsql;
```

Table Partitioning before PostgreSQL 10

- A trigger needs to be added to the parent table

```
CREATE TRIGGER partition_audit_trigger
    BEFORE INSERT ON audit
    FOR EACH ROW
    EXECUTE PROCEDURE
        partition_audit_trigger();
```

Execution Performance

- Performance is much better if dynamic SQL is not used

```
CREATE OR REPLACE FUNCTION partition_audit_trigger()
RETURNS TRIGGER AS $$

BEGIN
    IF ( NEW.event_time >= DATE '2015-01-01' AND
        NEW.event_time < DATE '2016-01-01' ) THEN
        INSERT INTO audit_2015 VALUES (NEW.*);
    ELSIF ( NEW.event_time >= DATE '2014-01-01' AND
        NEW.event_time < DATE '2015-01-01' ) THEN
        INSERT INTO audit_2014 VALUES (NEW.*);
    ELSE
        RAISE EXCEPTION 'Date out of range. Fix
                        partition_audit_trigger() function!';
    END IF;
    RETURN NULL;
END;
$$ LANGUAGE plpgsql;
```

Moving Partitions

- If the column used for the partition key changes, the row may need to be moved to a different partition

```
CREATE TRIGGER move_partition_audit_trigger
    BEFORE UPDATE
        ON audit_2014
    FOR EACH ROW EXECUTE PROCEDURE
        move_partition_audit_trigger('2014-01-01', '2015-01-01');
```

```
CREATE TRIGGER move_partition_audit_trigger
    BEFORE UPDATE
        ON audit_2015
    FOR EACH ROW EXECUTE PROCEDURE
        move_partition_audit_trigger('2015-01-01', '2016-01-01');
```

Moving Partitions

```
CREATE FUNCTION move_partition_audit_trigger() RETURNS TRIGGER AS $$  
DECLARE  
    start_date DATE;  
    end_date DATE;  
BEGIN  
    start_date := TG_ARGV[0];  
    end_date := TG_ARGV[1];  
    IF ( NEW.event_time IS DISTINCT FROM OLD.event_time ) THEN  
        IF (NEW.event_time < start_date OR NEW.event_time >= end_date) THEN  
            EXECUTE 'DELETE FROM ' || TG_TABLE_SCHEMA || '.' ||  
                    TG_TABLE_NAME || ' WHERE ctid = $1' USING OLD.ctid;  
            INSERT INTO audit VALUES (NEW.*);  
            RETURN null;  
        END IF;  
    END IF;  
    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;
```

Moving Partitions

- Only fire the trigger if the partition key changes

```
CREATE TRIGGER move_partition_audit_trigger
    BEFORE UPDATE
        ON audit_2014
    FOR EACH ROW
    WHEN (NEW.event_time IS DISTINCT FROM OLD.event_time)
        EXECUTE PROCEDURE
move_partition_audit_trigger('2014-01-01', '2015-01-01');
```

```
CREATE TRIGGER move_partition_audit_trigger
    BEFORE UPDATE
        ON audit_2015
    FOR EACH ROW
    WHEN (NEW.event_time IS DISTINCT FROM OLD.event_time)
        EXECUTE PROCEDURE
move_partition_audit_trigger('2015-01-01', '2016-01-01');
```

Trigger Use Cases

- Calculate columns
 - Calculate complex values
 - Extract values from complex structures
 - Enforce derived values when using denormalization
 - Used to:
 - Increase performance
 - Simplify queries

Extract JSON

```
$ head -n 2 zips.json
{ "_id" : "01001", "city" : "AGAWAM",
  "loc" : [ -72.622739, 42.070206 ], "pop" : 15338, "state" : "MA" }
{ "_id" : "01002", "city" : "CUSHMAN",
  "loc" : [ -72.51564999999999, 42.377017 ], "pop" : 36963, "state" :
"MA" }
CREATE TABLE zips (
    zip_code varchar PRIMARY KEY,
    state    varchar,
    data     json
);

```

Extract JSON

```
CREATE OR REPLACE FUNCTION extract_data_trigger()
    RETURNS TRIGGER AS $$

BEGIN
    NEW.zip_code := NEW.data->>'_id';
    NEW.state := NEW.data->>'state';

    RETURN NEW;
END;
$$
LANGUAGE plpgsql;

CREATE TRIGGER extract_data_trigger
    BEFORE UPDATE OR INSERT ON zips
    FOR EACH ROW EXECUTE PROCEDURE extract_data_trigger();
```

Trigger Use Cases

- Cache invalidation
 - Remove stale entries from a cache
 - The database tracks all data so is the single source of truth
 - Used to:
 - Simplify cache management
 - Remove application complexity

Note: Foreign Data Wrappers simplify this process significantly

Note: ON (action) CASCADE constraints can simplify this too.

Cache Invalidation

```
CREATE FUNCTION remove_cache_trigger()
    RETURNS TRIGGER AS $$

BEGIN
    DELETE from myredis_cache
    WHERE key = OLD.id::varchar;

    RETURN NEW;
END;
$$
LANGUAGE plpgsql;

CREATE TRIGGER remove_cache_trigger
AFTER UPDATE OR DELETE ON users
FOR EACH ROW EXECUTE PROCEDURE remove_cache_trigger();
```

Cache Invalidation - Async

- The latency of updating the cache may not be an acceptable as part of the main transaction

```
CREATE FUNCTION remove_cache_trigger()
    RETURNS TRIGGER AS $$

BEGIN
    PERFORM pg_notify(TG_TABLE_NAME, OLD.id::varchar);

    RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

Event Triggers

- Event triggers fire for DML commands (CREATE, ALTER, DROP, etc)
- They are not tied to a single table
- They are global to a database

Event Triggers

```
CREATE OR REPLACE FUNCTION notice_ddl()
    RETURNS event_trigger AS
$$
BEGIN
    RAISE NOTICE 'DDL Fired: % %', tg_event, tg_tag;
END;
$$
LANGUAGE plpgsql;
```

```
CREATE EVENT TRIGGER notice_ddl
    ON ddl_command_start
    EXECUTE FUNCTION notice_ddl();
```

Event Trigger Events

- `ddl_command_start`
- `ddl_command_end`
- `table_rewrite`
- `sql_drop`

ddl_command_start

- Fired just before the command starts
 - This is before any information is known about the command
- Fires for all event trigger command tags
- Does not fire for shared objects such as databases and roles
- Does not fire for commands involving event triggers

`ddl_command_end`

- Fired after the command ends
- Fires for all event trigger command tags
- The objects have been affected so the details of the command can be obtained

sql_drop

- Fired just before ddl_command_end fires
 - The objects have already been removed so they are not accessible
- Only fired for commands that drop an object
- The objects have been affected so the details of the command can be obtained

table_rewrite

- Fired just before the table is rewritten by the command
- Only fired for commands that rewrites an object
- CLUSTER and VACUUM FULL do not fire the event

Event Tags

- ALTER POLICY
- ALTER SCHEMA
- ALTER SEQUENCE
- ALTER TABLE
- CREATE EXTENSION
- CREATE FUNCTION
- CREATE INDEX
- CREATE SEQUENCE
- CREATE TABLE
- CREATE TABLE AS
- CREATE VIEW
- DROP INDEX
- DROP TABLE
- DROP VIEW
- GRANT
- REVOKE

The full list is available in the documentation

<https://www.postgresql.org/docs/current/event-trigger-matrix.html>

Event Trigger Functions

A set of functions to help retrieve information from event triggers

- pg_event_trigger_ddl_commands
- pg_event_trigger_dropped_objects
- pg_event_trigger_table_rewrite_oid
- pg_event_trigger_table_rewrite_reason

Understanding pg_event_trigger_ddl_commands

- Returns a line for each DDL command executed
- Only valid inside a ddl_command_end trigger

Column	Type
classid	oid
objid	oid
objsubid	integer
command_tag	text
object_type	text
schema_name	text
object_identity	text
in_extension	bool
command	pg_ddl_command

Understanding pg_event_trigger_dropped_objects

- Returns a line for each object dropped by the DDL executed
- Only valid inside a sql_drop trigger

Column	Type
classid	oid
objid	oid
objsubid	integer
original	bool
normal	bool
is_temporary	bool

Column	Type
object_type	text
schema_name	text
object_name	text
object_identity	text
address_names	text[]
address_args	text[]

Understanding rewrite functions

- Only valid inside a `table_rewrite` trigger
- `pg_event_trigger_table_rewrite_oid`
 - Returns the OID of the table about to be rewritten
- `pg_event_trigger_table_rewrite_reason`
 - Returns the reason code of why the table was rewritten

Using Event Trigger Functions

```
CREATE OR REPLACE FUNCTION stop_drops()
    RETURNS event_trigger AS
$$
DECLARE
    l_tables varchar[] := '{sales, inventory}';
BEGIN
    IF EXISTS(SELECT 1
              FROM pg_event_trigger_dropped_objects()
              WHERE object_name = ANY (l_tables)) THEN
        RAISE EXCEPTION 'Drops of critical tables are not permitted';
    END IF;
END;
$$ LANGUAGE plpgsql;

CREATE EVENT TRIGGER stop_drops
    ON sql_drop
    EXECUTE FUNCTION stop_drops();
```

Things to Remember

- Triggers impose a performance overhead on transactions
- Triggers are part of the parent transaction
 - The trigger fails, the main transaction fails
 - The main transaction rolls back, the trigger call never happened
 - If the trigger takes a long time, the whole transaction timing is affected
- Triggers can be difficult to debug
 - Especially cascaded triggers

PL/pgSQL Best Practices



Programming Practices

- Follow good programming practices
 - Indent code consistently
 - Comment code liberally
 - Code reuse/modularity practices are different than other programming languages
 - Deep call stacks in PL/pgSQL can be performance intensive

Naming Conventions

- Create and follow a consistent naming convention for objects
 - PostgreSQL is case insensitive so init cap does not work, use "_" to separate words in names
 - Prefix all parameter names with something like "p_"
 - Prefix all variable names with something like "v_"

Performance

- Avoid expensive constructs unless necessary
 - Dynamic SQL
 - EXCEPTION blocks