



Database Migration Workshop

AWS Database Migration Service (DMS)

Oracle to Aurora PostgreSQL

July 2019

Table of Contents

Overview	3
Connecting to your Environment	4
Database Migration Service (DMS).....	5
Log on the EC2 instance.....	5
Connect to the Source Oracle on Amazon RDS	6
Create Replication Instance	15
Create Source and Target Endpoints.....	19
Create and Run Your Database Migration Task.....	23
Inspect the Content of Target Database.....	28
Capture and Replicate Data Changes.....	29
Replicating Changes from Source to the Target Database	33
Summary.....	36

Overview

AWS Database Migration Service (DMS) helps you migrate databases to AWS easily and securely. The source database remains fully operational during the migration, minimizing downtime to applications that rely on the database. AWS DMS can migrate your data to and from most widely used commercial and open-source databases. The service supports homogenous migrations such as Oracle Server to Oracle Server or PostgreSQL to Aurora PostgreSQL, as well as heterogeneous migrations between different database platforms, such as SQL Server to Amazon Aurora or Oracle to PostgreSQL. AWS DMS can also be used for continuous data replication with high-availability.

In this lab session we will perform database migration from a source Oracle Server on an Amazon RDS instance to a target Amazon Aurora (PostgreSQL) database.

This lab will walk you through:

- Performing a full load migration of source oracle server to target Aurora PostgreSQL database using AWS DMS.
- Capturing change data set (CDC) from the Oracle Server EC2 instance to an RDS Aurora PostgreSQL instance using AWS DMS.



Connecting to your Environment

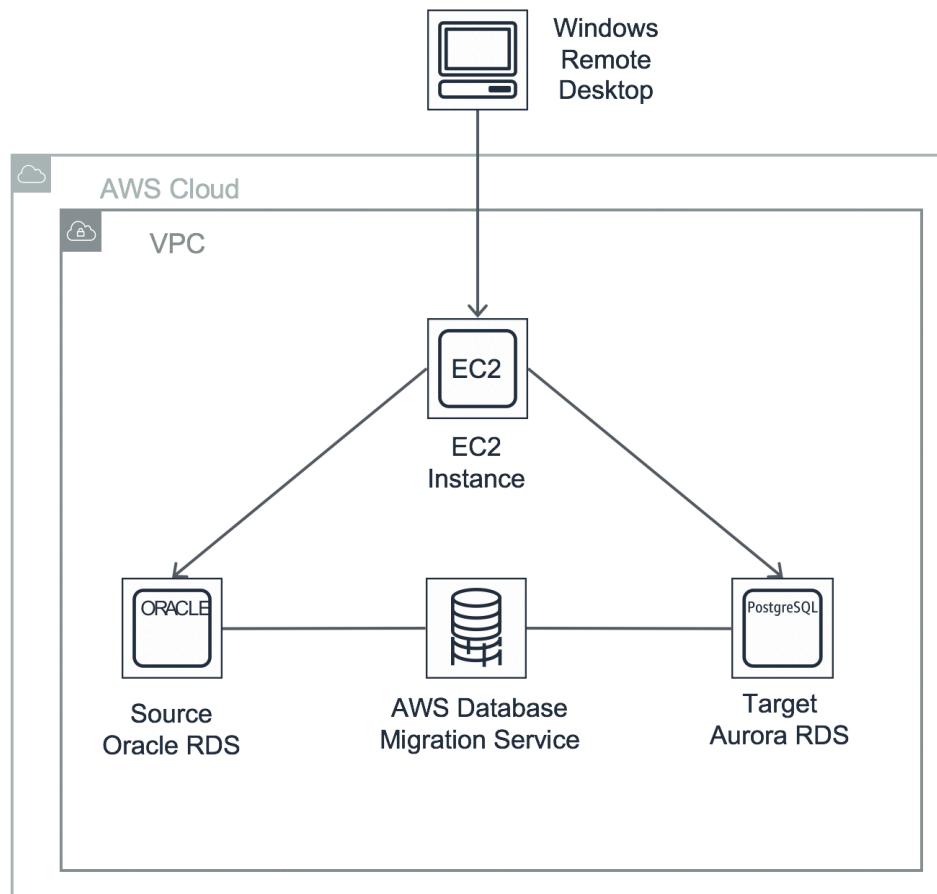
Before launching this lab, make sure you have completed the instructions in the following guide to setup your lab environment: <https://bit.ly/2xL3SMw> (Document '0-EnvironmentConfiguration.pdf')

The environment for this lab consists of:

- An EC2 instance used to run the AWS Schema Conversion Tool (SCT) as well as other applications needed to complete the lab.
- An Amazon RDS instance used to host the source Oracle database
- An Amazon RDS Aurora PostgreSQL instance used as the target database

Once you have completed the instructions in the above referenced document, take special note of the following output values:

- **SourceEC2PublicDNS**
- **SourceOracleEndpoint**
- **TargetAuroraPostgreSQLEndpoint**



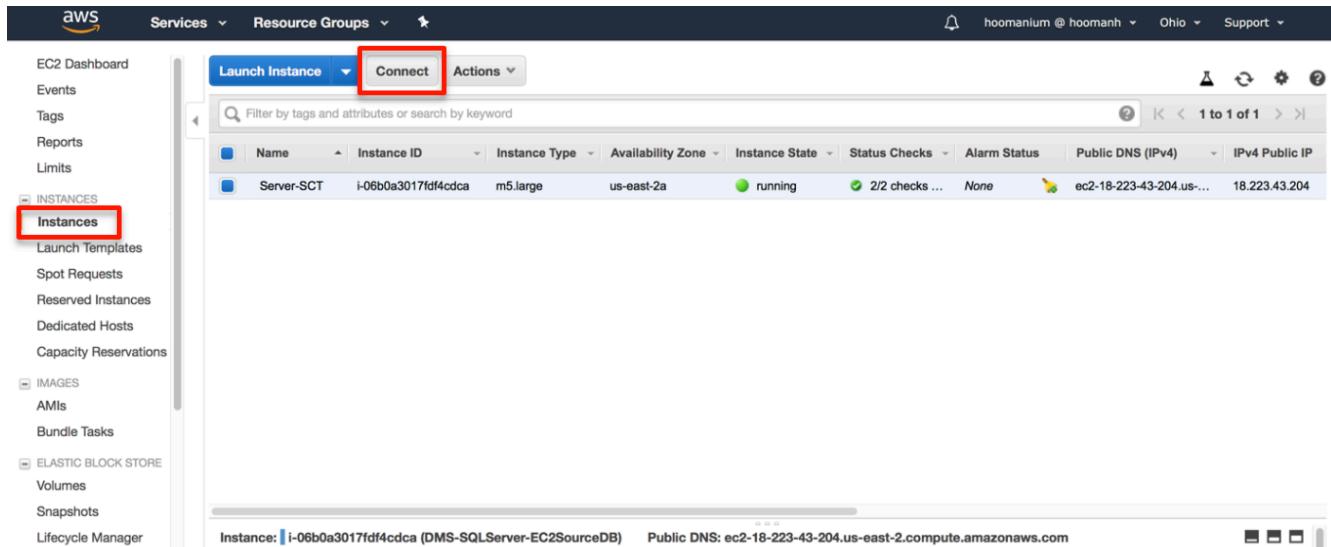
Database Migration Service (DMS)

The following steps provide instructions to migrate existing data from the source Oracle running on an Amazon RDS instance to a target Amazon Aurora (PostgreSQL) database. In this exercise you perform the following tasks:

- Connect the source Oracle database
- Configure the source database for replication
- Configure the target database for replication
- Create an AWS DMS Replication Instance
- Create AWS DMS source and target endpoints
- Create and run your AWS DMS migration task
- Replicate ongoing changes from the source to the target database

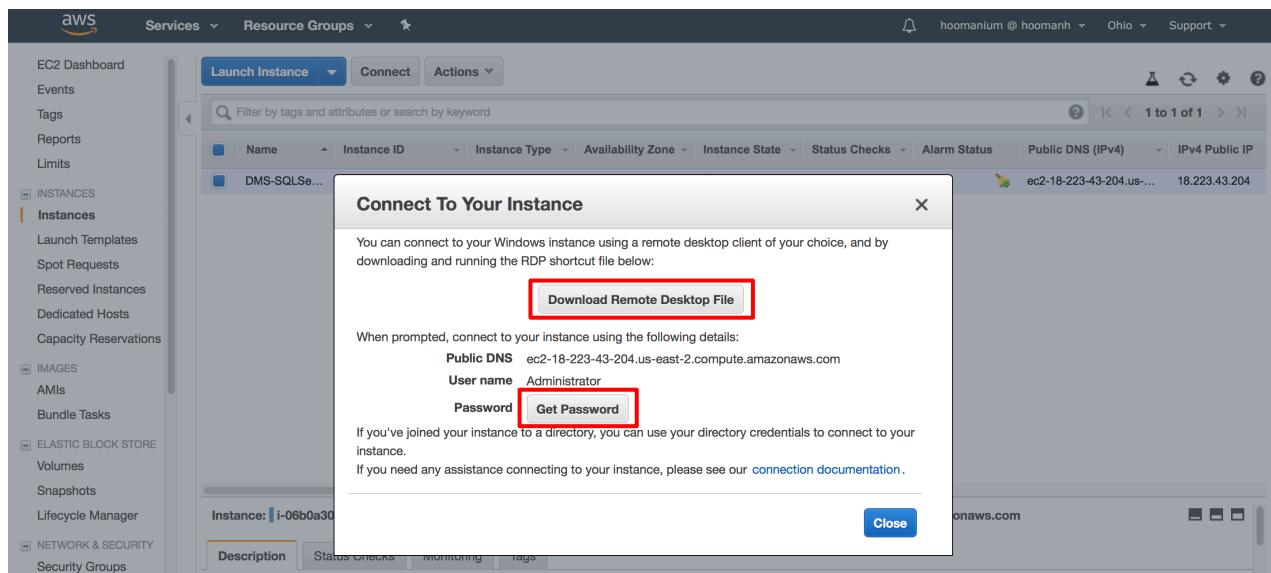
Log on the EC2 instance

1. Go to the AWS EC2 [REDACTED] and click on **Instances** in the left column.
2. Select the instance with the name **Server-SCT** and then click the **Connect** button.



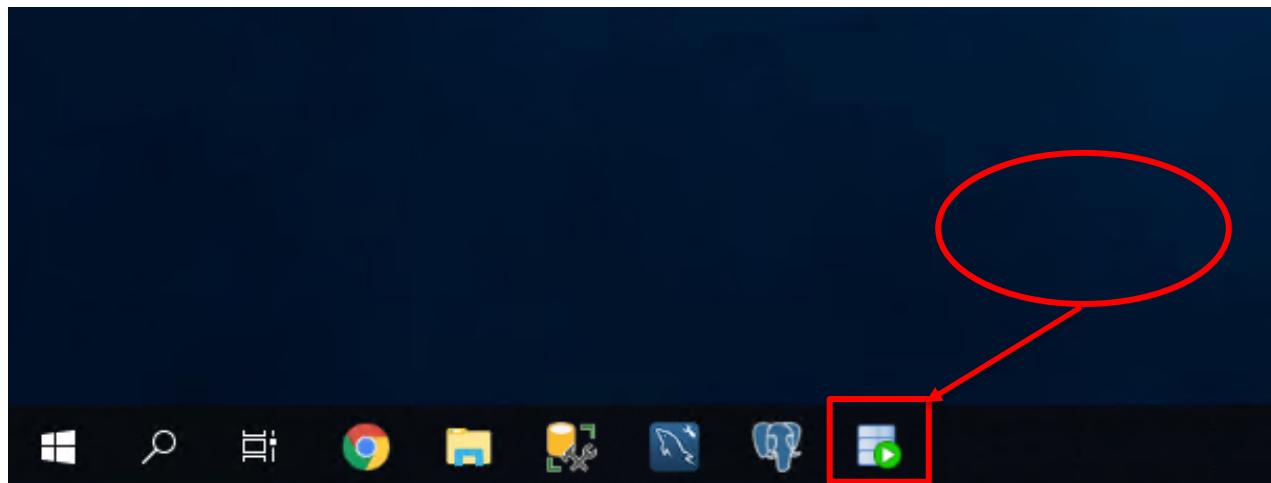
3. In this step, you perform 3 tasks:

- Click the **Get Password** button and upload the **Key Pair** file that you downloaded earlier. Please take note of the EC2 console generated administrator password.
- Click on **Download Remote Desktop File** to download the RDP file to access this EC2 instance.
- Connect to the EC2 instance using the RDP.



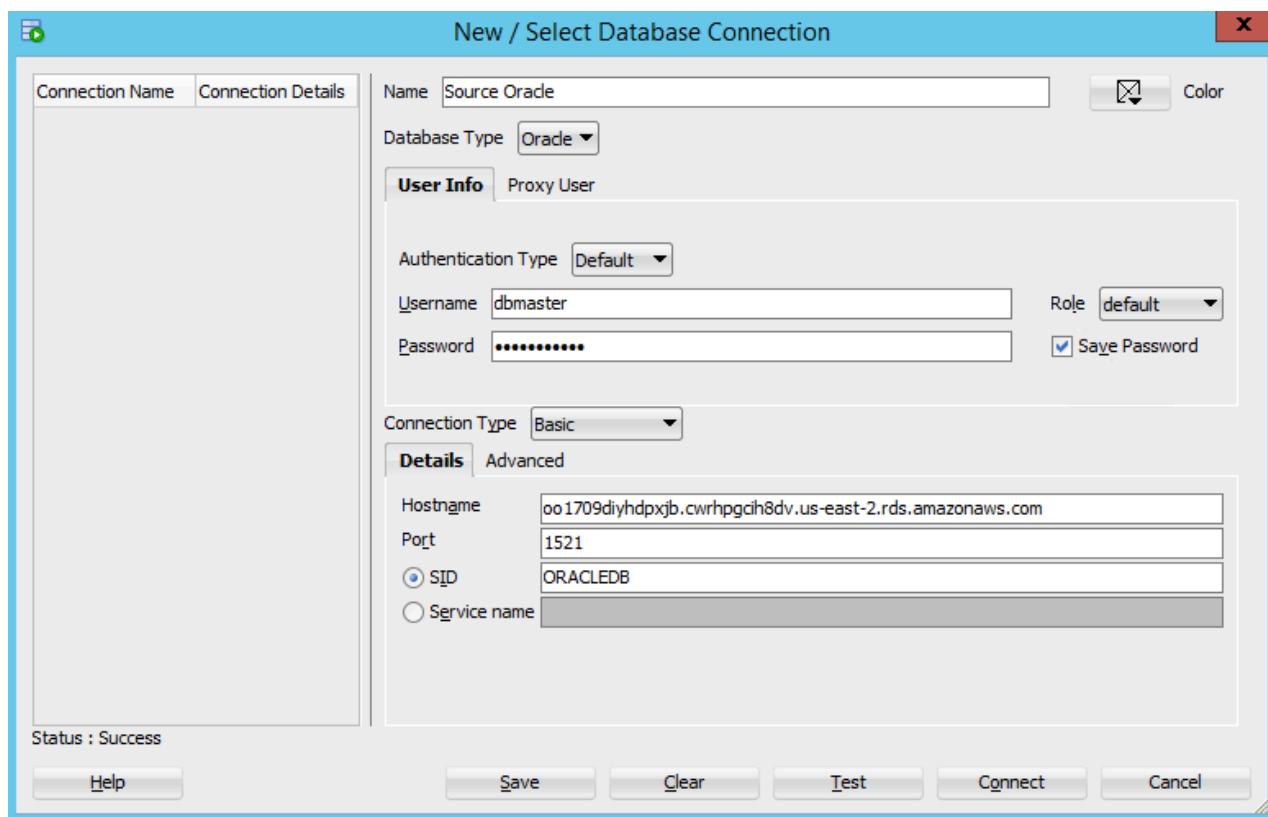
Connect to the Source Oracle on Amazon RDS

4. Once you have connected to the EC2 instance, open **Oracle SQL Developer** from the Taskbar.



5. Click on the **plus sign** from the left-hand menu to create a **New Database Connection** using the following values, then click **Connect**.

Parameter	Value
Connection Name	Source Oracle
Username	dbmaster
Password	dbmaster123
Save Password	Check
Hostname	< SourceOracleEndpoint >
Port	1521
SID	ORACLEDDB



6. After the you see the test status as **Successful**, click **Connect**.

Configure the Source Oracle database for Replication

To use Oracle as a source for AWS Database Migration Service (AWS DMS), you must first provide a user account (DMS user) with read and write privileges on the Oracle database.

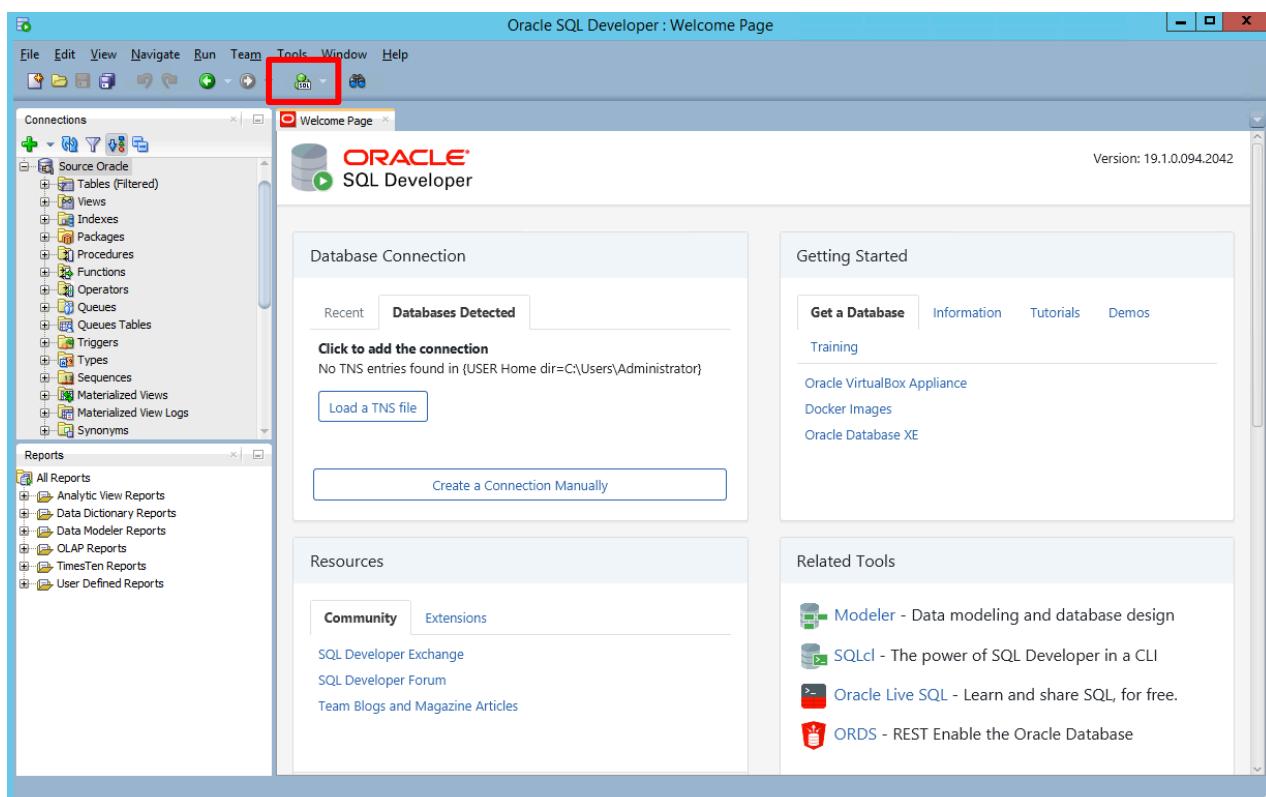
You also need to ensure that ARCHIVELOG MODE is on to provide information to LogMiner. AWS DMS uses LogMiner to read information from the archive logs so that AWS DMS can capture changes.

For AWS DMS to read this information, make sure the archive logs are retained on the database server as long as AWS DMS requires them. Retaining archive logs for 24 hours is usually sufficient.

To capture change data, AWS DMS requires database-level supplemental logging to be enabled on your source database. Doing this ensures that the LogMiner has the minimal information to support various table structures such as clustered and index-organized tables.

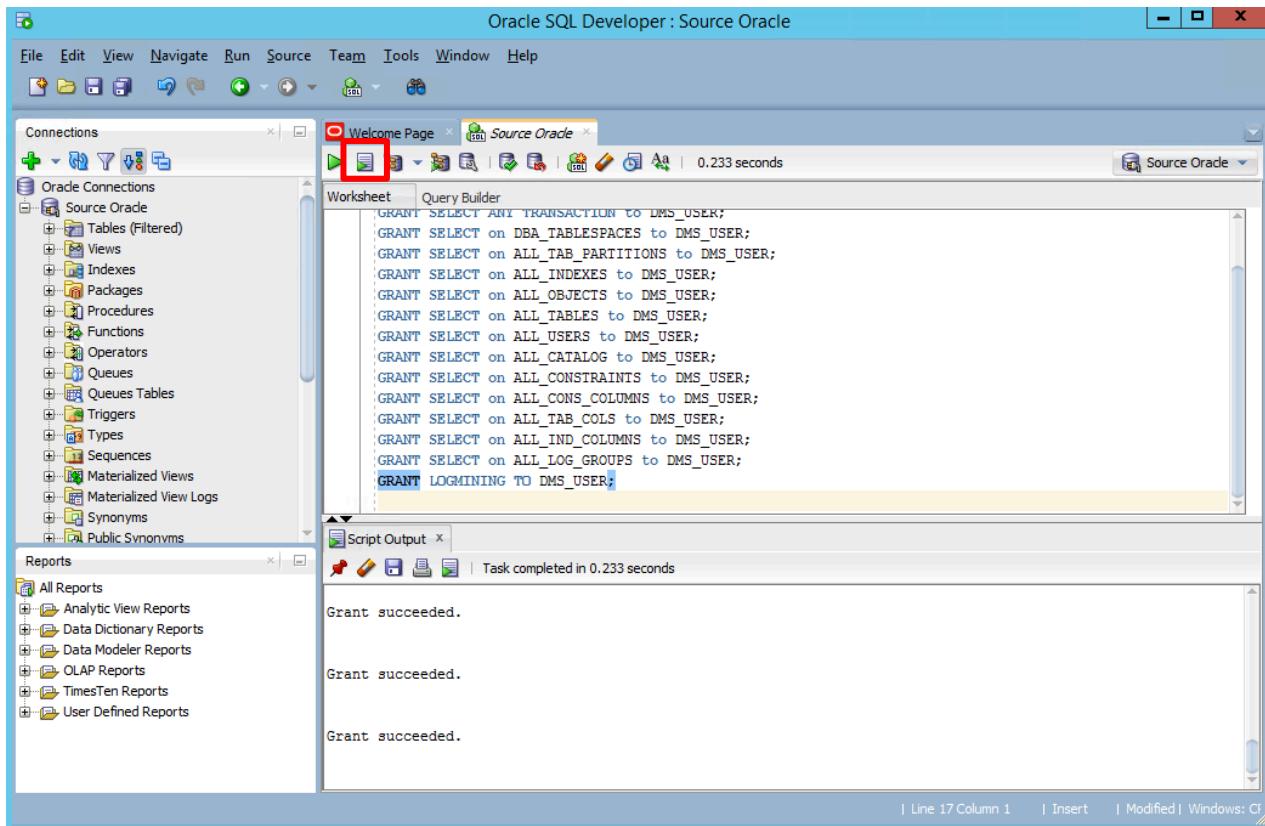
Similarly, you need to enable table-level supplemental logging for each table that you want to migrate.

7. Click on the **SQL Worksheet** icon within **Oracle SQL Developer**, then connect to the **Source Oracle** database.



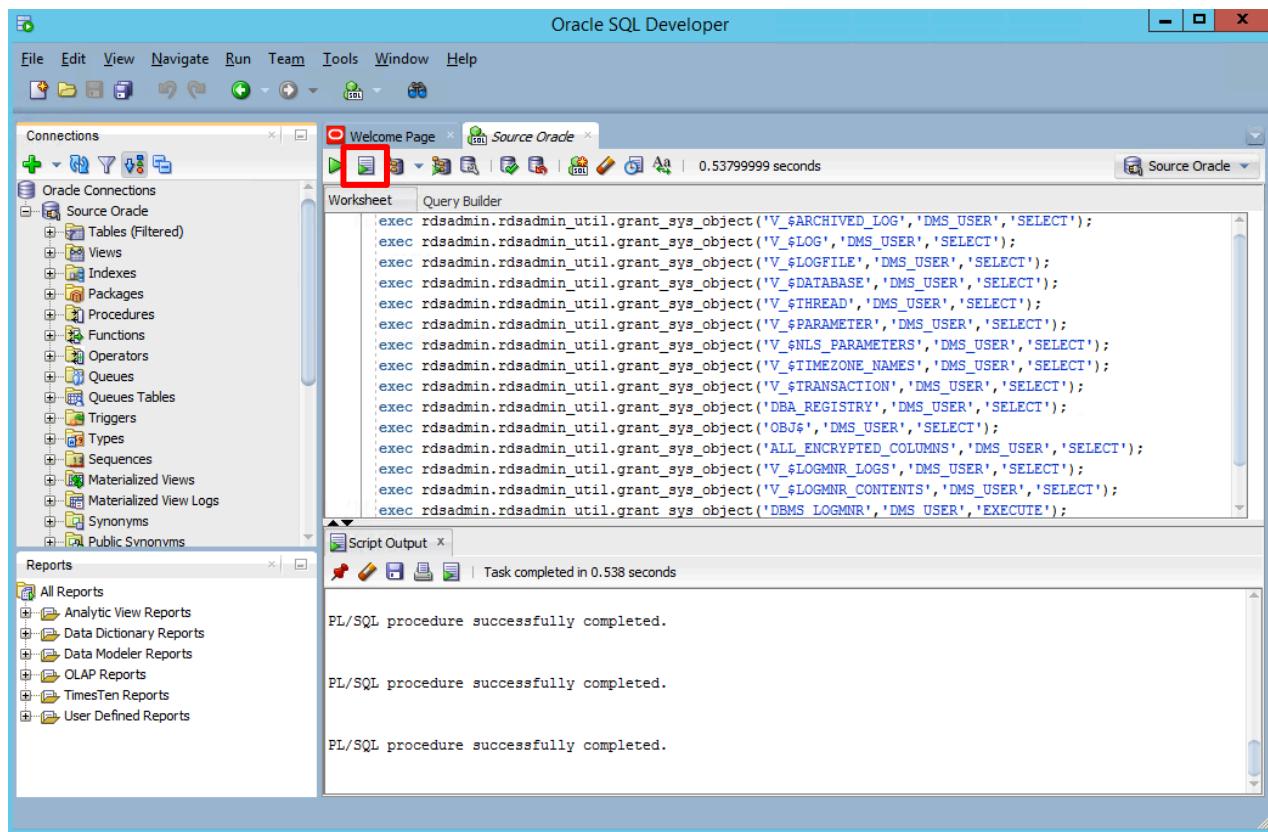
8. Next, execute the below statements to grant the following privileges to the AWS DMS user to access the source Oracle endpoint:

```
GRANT SELECT ANY TABLE to DMS_USER;
GRANT SELECT on ALL_VIEWS to DMS_USER;
GRANT SELECT ANY TRANSACTION to DMS_USER;
GRANT SELECT on DBA_TABLESPACES to DMS_USER;
GRANT SELECT on ALL_TAB_PARTITIONS to DMS_USER;
GRANT SELECT on ALL_INDEXES to DMS_USER;
GRANT SELECT on ALL_OBJECTS to DMS_USER;
GRANT SELECT on ALL_TABLES to DMS_USER;
GRANT SELECT on ALL_USERS to DMS_USER;
GRANT SELECT on ALL_CATALOG to DMS_USER;
GRANT SELECT on ALL_CONSTRAINTS to DMS_USER;
GRANT SELECT on ALL_CONS_COLUMNS to DMS_USER;
GRANT SELECT on ALL_TAB_COLS to DMS_USER;
GRANT SELECT on ALL_IND_COLUMNS to DMS_USER;
GRANT SELECT on ALL_LOG_GROUPS to DMS_USER;
GRANT LOGMINING TO DMS_USER;
```



9. In addition, run the following:

```
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$ARCHIVED_LOG','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$LOG','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$LOGFILE','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$DATABASE','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$THREAD','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$PARAMETER','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$NLS_PARAMETERS','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$TIMEZONE_NAMES','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$TRANSACTION','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('DBA_REGISTRY','DMS_USER','SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('OBJ$', 'DMS_USER', 'SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('ALL_ENCRYPTED_COLUMNS', 'DMS_USER', 'SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$LOGMNR_LOGS', 'DMS_USER', 'SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('V_$LOGMNR_CONTENTS', 'DMS_USER', 'SELECT');
exec rdsadmin.rdsadmin_util.grant_sys_object('DBMS_LOGMNR', 'DMS_USER', 'EXECUTE');
```



10. Run the following query to retain archived redo logs of the source Oracle database instance for 24 hours:

```
exec rdsadmin.rdsadmin_util.set_configuration('archivelog retention hours',24);
```

11. Run the following query to enable database-level supplemental logging:

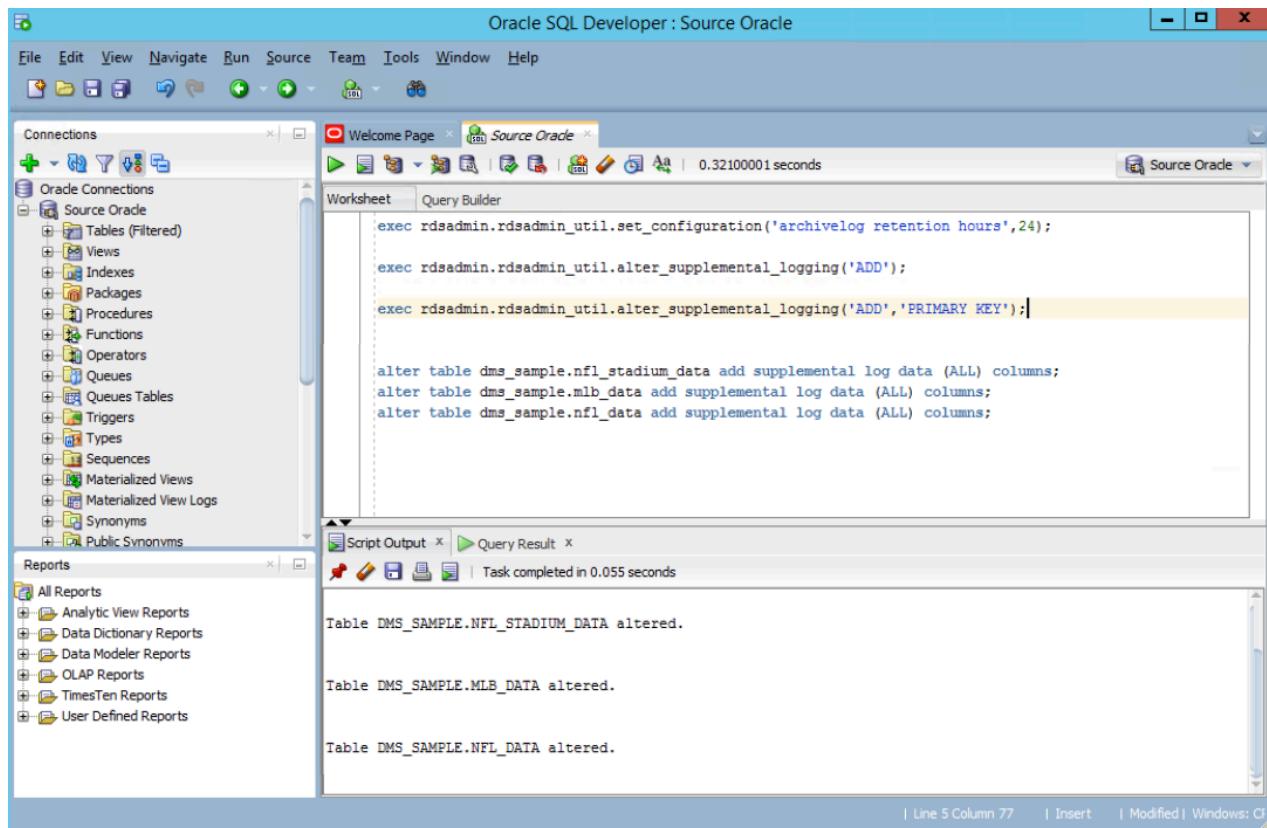
```
exec rdsadmin.rdsadmin_util.alter_supplemental_logging('ADD');
```

12. Run the following query to enable PRIMARY KEY logging for tables that have primary keys:

```
exec rdsadmin.rdsadmin_util.alter_supplemental_logging('ADD','PRIMARY KEY');
```

13. Run the following queries to add supplemental logging for tables that don't have primary keys, use the following command to add supplemental logging:

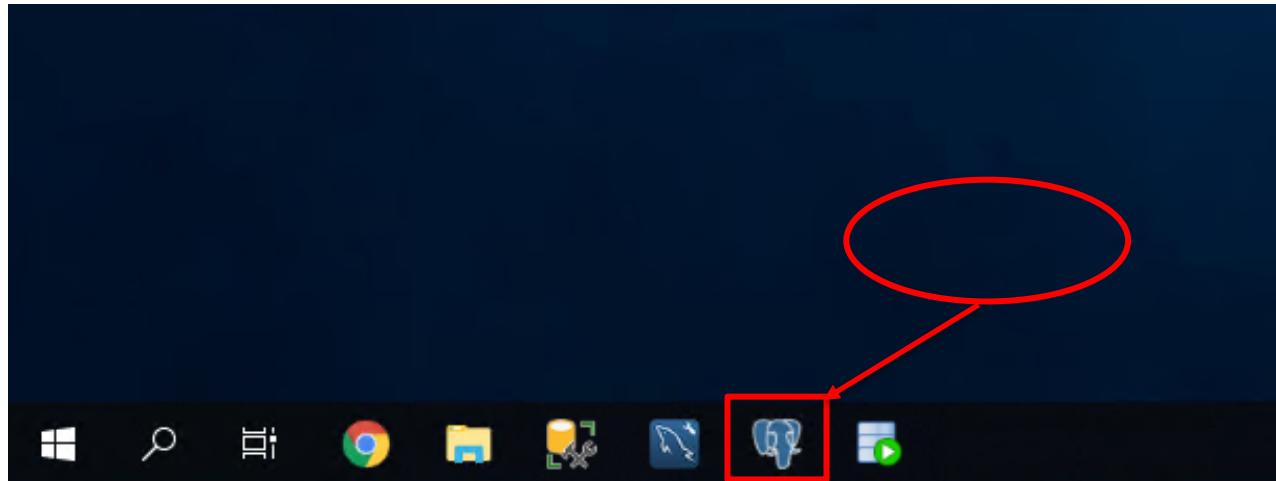
```
alter table dms_sample.nfl_stadium_data add supplemental log data (ALL) columns;
alter table dms_sample.mlb_data add supplemental log data (ALL) columns;
alter table dms_sample.nfl_data add supplemental log data (ALL) columns;
```



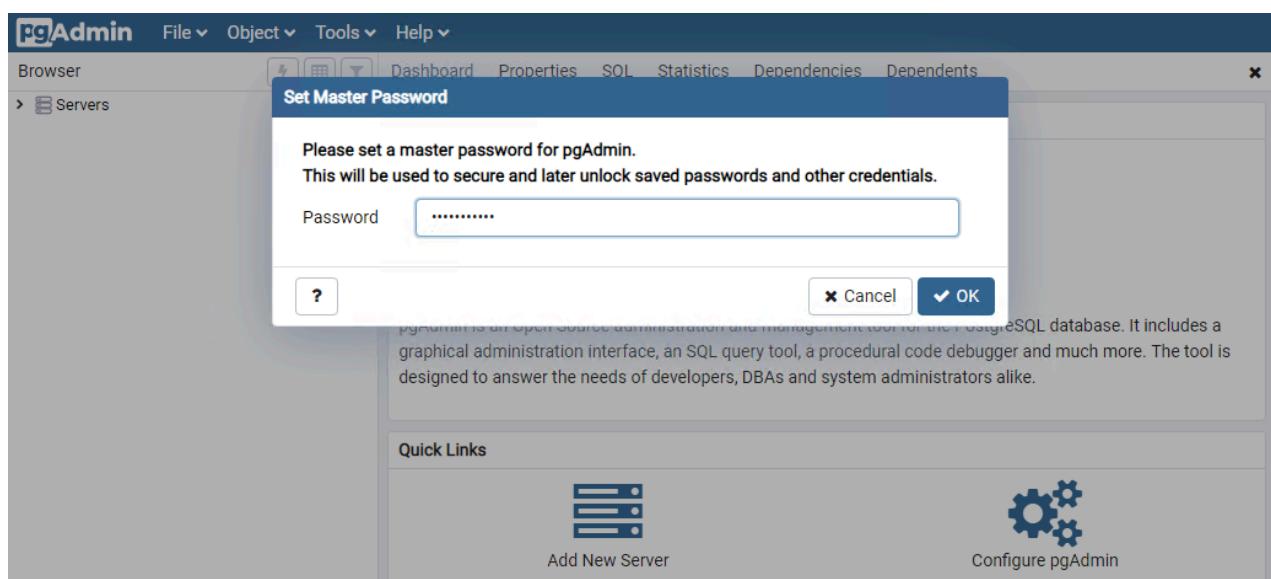
Configure the Target Aurora RDS database for Replication

During the full load process, AWS DMS does not load tables in any particular order, so it might load the child table data before parent table data. As a result, foreign key constraints might be violated if they are enabled. Also, if triggers are present on the target database, they might change data loaded by AWS DMS in unexpected ways.

14. Open pgAdmin 4 from the **Taskbar** on the EC2 server.



15. You may be prompted to set a **Master Password**. Enter **dbmaster123**, then click, **OK**.



16. Click on the **Add New Server** icon, and enter the following values. Then, press **Save**.

Parameter	Value
General -> Name	Target Aurora RDS (PostgreSQL)
Connection -> Host Name/Address	< TargetAuroraPostgreSQLEndpoint >
Connection -> Port	5432
Connection -> Username	dbmaster
Connection -> Password	dbmaster123
Connection -> Save Password	Check

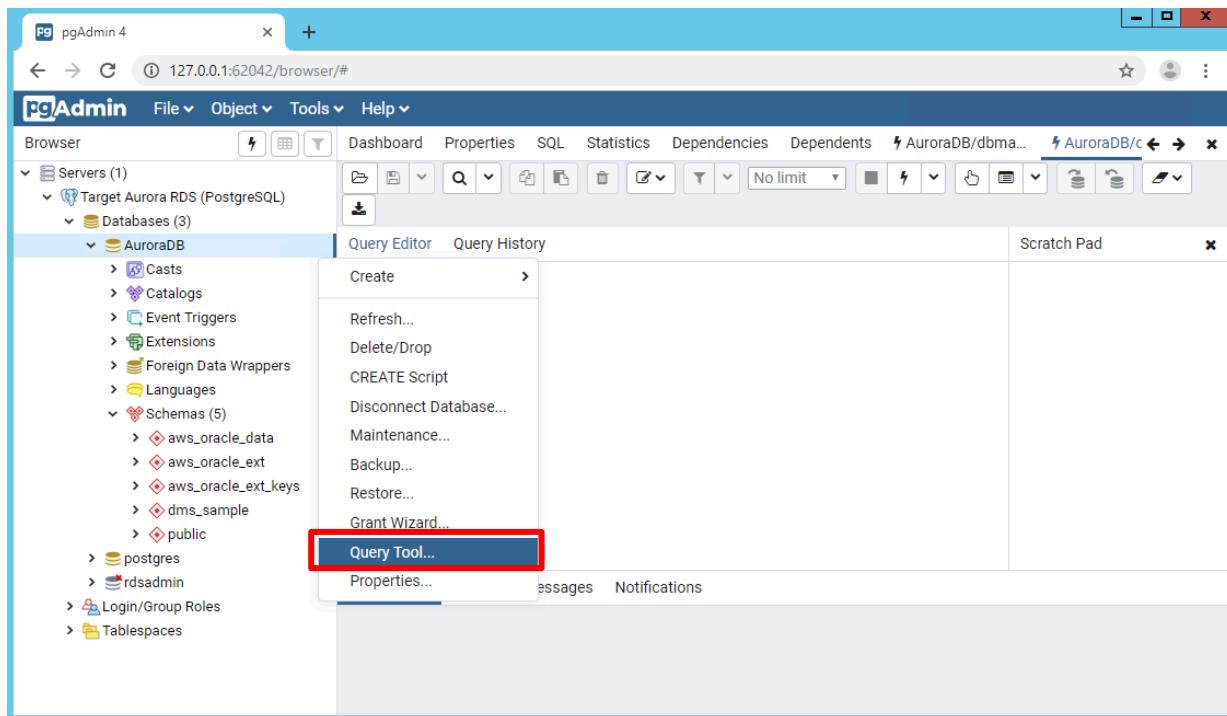
Create - Server

General Connection SSL SSH Tunnel Advanced

Host name/address	oracle-lab-auroracluster-1t2kfyq1nvy4.cluster-cwrhpg
Port	5432
Maintenance database	postgres
Username	dbmaster
Password
Save password?	<input checked="" type="checkbox"/>
Role	
Service	

Cancel **Reset** **Save**

17. Right-click on **AuroraDB** database from left-hand menu, and then select **Query Tools**.



18. In this step you are going to drop the foreign key constraints from the target database:

- Open **DropConstraintsPostgreSQL.sql** inside the **DMS Workshop\Scripts** folder in Notepad.
- Copy the content of the file to the **Query Editor** in pgAdmin 4.
- Execute** the script.

The screenshot shows the pgAdmin 4 interface with the 'Query Editor' tab active. The 'Execute' button in the toolbar is highlighted with a red box. The query editor contains the following SQL script:

```

27
28 ALTER TABLE dms_sample.sporting_event_ticket
29 DROP CONSTRAINT IF EXISTS set_sporting_event_fk;
30
31 ALTER TABLE dms_sample.sport_division
32 DROP CONSTRAINT IF EXISTS sd_sport_league_fk;
33
34 ALTER TABLE dms_sample.sport_division
35 DROP CONSTRAINT IF EXISTS sd_sport_type_fk;
36
37 ALTER TABLE dms_sample.sport_league
38 DROP CONSTRAINT IF EXISTS sl_sport_type_fk;

```

The results pane shows the output of the query:

```

Data Output Explain Messages Notifications
NOTICE: constraint "se_location_id_fk" of relation "sporting_event" does not exist, skipping
ALTER TABLE
Query returned successfully in 64 msec.

```

A green message box at the bottom right indicates: "Query returned successfully in 64 msec."

Create Replication Instance

The following illustration shows a high-level view of the migration process.



An AWS DMS replication instance performs the actual data migration between source and target. The replication instance also caches the transaction logs during the migration. The amount of CPU and memory capacity a replication instance has influences the overall time that is required for the migration.

19. Click on <https://console.aws.amazon.com/dms/> to launch the Database Migration Service.
20. On the left-hand menu click on **Replication Instances**. This will launch the Replication instance screen.
21. Click on the **Create replication instance** button on the top right side.

Name	Class	Status	Engine version	Availability zone	VPC	Public	Public IP address	Private
Empty replication instance table You don't have any replication instances.								

22. Enter the following information for the **Replication Instance**. Then, click on the **Create** button.

Parameter	Value
Name	oracle-replication
Description	Oracle to Aurora DMS replication instance
Instance Class	dms.c4.2xlarge This is not the default, please change it
Engine version	Leave the default value
Allocated storage (GB)	50
VPC	< VPC ID from Environment Setup Step >
Multi-AZ	No
Publicly accessible	No
Advanced -> VPC Security Group(s)	default

AWS Database Migration Lab – Oracle Migration

Screenshot of the AWS Database Migration Service (DMS) "Create replication instance" configuration page.

Replication instance configuration

Name: oracle-replication

Description: Oracle to Aurora DMS replication instance

Instance class: dms.c4.4xlarge

Billing: Billing is based on [DMS pricing](#).

Engine version: 3.1.3

Allocated storage (GB): 50

VPC: Choose an Amazon Virtual Private Cloud (VPC) where your replication instance should run. Selected: `vpc-0070bce6424a3e253 - DMS Oracle Lab`

Multi AZ: If you choose this option, AWS DMS will perform a multi-AZ deployment, with a primary instance in one availability zone (AZ) and a standby instance in another AZ. This configuration provides a highly available, fault-tolerant replication environment.

Billing: Billing is based on [DMS pricing](#).

Publicly accessible: If you choose this option, AWS DMS will assign a public IP address to your replication instance, and you'll be able to connect to databases outside of your Amazon VPC.

Advanced security and network configuration

Replication subnet group: default-vpc-0070bce6424a3e253

Availability zone: No Preference

VPC security group(s): Oracle-Lab-OraVPCSecurityGroup-TT003UN10AKS

KMS master key: (Default) aws/dms

Maintenance:

Cancel **Create**

NOTE: Creating replication instance will take several minutes. While waiting for the replication instance to be created, you can specify the source and target database endpoints in the next steps. However, you can test connectivity only after the replication instance has been created, because the replication instance is used in the connection.

Create Source and Target Endpoints

Now that you have a replication instance, you need to create source and target endpoints for the sample database.

- Click on the **Endpoints** link on the left, and then click on **Create endpoint** on the top right corner.

The screenshot shows the AWS DMS console. On the left, there's a sidebar with 'AWS DMS' at the top, followed by 'Dashboard', 'Conversion & migration' (with 'Database migration tasks' under it), 'Resource management' (with 'Replication instances' under it), and 'Endpoints' (which is highlighted with a red box). The main area is titled 'Endpoint' and shows a table with one row: 'Empty endpoint table' and 'You don't have any endpoints.' At the top right of the main area, there's a 'Create endpoint' button, which is also highlighted with a red box.

- Enter the following information to create an endpoint for the source **dms_sample** database (continued on the next page):

Parameter	Value
Endpoint Type	Source endpoint
Select RDS DB instance	Check
RDS Instance	<StackName>-SourceOracleDB
Endpoint Identifier	oracle-source
Source Engine	oracle
Server Name	< SourceOracleEndpoint >
Port	1521
SSL Mode	none
User Name	dbmaster
Password	dbmaster123
SID/Service Namee	ORACLEDDB
Test endpoint connection -> VPC	< VPC ID from Environment Setup Step >
Replication Instance	oracle-replication

AWS Database Migration Lab – Oracle Migration

Screenshot of the AWS Database Migration Service (DMS) "Create endpoint" wizard.

Endpoint type: Source endpoint (selected).
Target endpoint (unselected).

Select RDS DB instance: Select RDS DB instance.
RDS Instance: oo1709diyhdpxjb (dropdown menu).

Endpoint configuration:

- Endpoint identifier:** Oracle-Source.
- Source engine:** oracle.
- Server name:** oo1709diyhdpxjb.cwrrhpgch8dv.us-east-2.rds.amazonaws.com
- Port:** 1521.
- Secure Socket Layer (SSL) mode:** none.
- User name:** dbmaster.
- Password:** (empty field).
- SID/Service name:** ORACLEDDB.

Endpoint-specific settings: (Expander panel)

KMS master key: (Expander panel)

Test endpoint connection (optional): (Expander panel)

Test your endpoint connection by selecting a replication instance within your desired VPC. After clicking "Run test", an endpoint will be created with the details provided and attempt to connect to the instance. If the connection fails, you can edit and test it again. Endpoints that aren't saved will be deleted.

VPC: vpc-0070bce6424a3e253 - DMS Oracle Lab.

Replication instance: oracle-replication.

Run test: (button)

After clicking "Run test", an endpoint will be created with the details provided and attempt to connect to the instance. If the connection fails, you can edit and test it again. Endpoints that aren't saved will be deleted.

Endpoint identifier	Replication instance	Status	Message
No records found			

Create endpoint (button)

25. Once the information has been entered, click **Run Test**. When the status turns to **successful**, click **Create endpoint**.
26. Follow the same steps to create another endpoint for the **Target Aurora RDS Database** using the following values:

Parameter	Value
Endpoint Type	Target endpoint
Select RDS DB instance	Check
RDS Instance	<StackName>-AuroraPostgreSQLInstance
Endpoint Identifier	aurora-target
Source Engine	aurora-postgresql
Server Name	< TargetAuroraPostgreSQLEndpoint >
Port	1521 5432
SSL Mode	none
User Name	dbmaster
Password	dbmaster123
Database Name	AuroraDB
Test endpoint connection -> VPC	< VPC ID from Environment Setup Step >
Replication Instance	oracle-replication

27. Once the information has been entered, click **Run Test**. When the status turns to **successful**, click **Create endpoint**.

AWS Database Migration Lab – Oracle Migration

The screenshot shows the AWS DMS 'Create endpoint' wizard. The left sidebar lists navigation options like Dashboard, Conversion & migration, Resource management, Endpoints (selected), Certificates, Subnet groups, Events, Event subscriptions, and What's new.

Endpoint type: Target endpoint (selected). A source endpoint allows AWS DMS to read data from a database (on-premises or in the cloud), or from other data source such as Amazon S3.

Select RDS DB instance: checked. RDS Instance dropdown: oa1a10md2w8566r.cwrhpgcih8dv.us-east-2.rds.amazonaws.com

Endpoint configuration:

- Endpoint identifier:** aurora-Target
- Target engine:** aurora-postgresql
- Server name:** oa1a10md2w8566r.cwrhpgcih8dv.us-east-2.rds.amazonaws.com
- Port:** 5432
- Secure Socket Layer (SSL) mode:** none
- User name:** dbmaster
- Password:** (empty)
- Database name:** AuroraDB

Endpoint-specific settings: (button)

KMS master key: (button)

Test endpoint connection (optional):

Test your endpoint connection by selecting a replication instance within your desired VPC. After clicking "Run test", an endpoint will be created with the details provided and attempt to connect to the instance. If the connection fails, you can edit and test it again. Endpoints that aren't saved will be deleted.

VPC: vpc-0070bce6424a3e253 - DMS Oracle Lab

Replication instance: oracle-replication

Run test: (button)

After clicking "Run test", an endpoint will be created with the details provided and attempt to connect to the instance. If the connection fails, you can edit and test it again. Endpoints that aren't saved will be deleted.

Endpoint identifier	Replication instance	Status	Message
No records found			

Create endpoint (button)

Create and Run Your Database Migration Task

AWS DMS uses **Database Migration Task** to migrate the data from source to the target database. In this part of the lab you are going to create two Database Migration Tasks: one for migrating the existing data, and another for capturing data changes on the source database and replicating the changes to the target database.

28. Click on **Database migration tasks** on the left-hand menu, then click on the **Create task** button on the top right corner.

The screenshot shows the AWS DMS console interface. On the left, there's a navigation sidebar with options like 'Dashboard', 'Conversion & migration' (which is expanded and has 'Database migration tasks' selected), and 'Resource management'. The main area is titled 'Database migration tasks' and shows a table with one row: 'Empty replication task table' and 'You don't have any replication tasks.' At the top right of this area, there's a 'Create task' button, which is also highlighted with a red box.

29. Create a data migration task with the following values for migrating the **dms_sample** database.

Parameter	Value
Task identifier	oracle-migration-task
Replication instance	oracle-replication
Source database endpoint	oracle-source
Target database endpoint	aurora-target
Migration type	Migrate existing data
Start task on create	Checked
Target table preparation mode	Do nothing
Include LOB columns in replication	Limited LOB mode
Max LOB size (KB)	32
Enable validation	Unchecked
Enable CloudWatch logs	Checked

30. Expand the **Table mappings** section, and select **Guided UI** for the editing mode

31. Click on **Add new selection rule** button and enter the following values:

Parameter	Value
Schema	DMS_SAMPLE
Table name	%
Action	Include

NOTE: If the Create Task screen does not recognize any schemas, make sure to go back to endpoints screen and click on your endpoint. Scroll to the bottom of the page and click on **Refresh Button (C)** in the **Schemas** section.

If your schemas still do not show up on the Create Task screen, click on the Guided tab and manually select **DMS_SAMPLE** schema and all tables.

32. Next, expand the **Transformation rules** section, and click on **Add new transformation rule**. Then, create the following rules:

Parameter	Value
Target	Schema
Schema Name	DMS_SAMPLE
Action	Make lowercase

Parameter	Value
Target	Table
Schema Name	DMS_SAMPLE
Table Name	%
Action	Make lowercase

Parameter	Value
Target	Column
Schema Name	DMS_SAMPLE
Table Name	%
Column Name	%
Action	Make lowercase

AWS Database Migration Lab – Oracle Migration

The screenshot shows the AWS DMS console with the 'Create replication task' wizard open. The left sidebar shows navigation links for Dashboard, Conversion & migration (with 'Database migration tasks' selected), Resource management, and What's new.

Task configuration:

- Task identifier: oracle-migration-task
- Replication instance: oracle-replication - vpc-0070bce6424a3e253
- Source database endpoint: oracle-source-correct
- Target database endpoint: aurora-target
- Migration type: Migrate existing data
- Start task on create

Task settings:

- Target table preparation mode: Truncate
- Include LOB columns in replication: Limited LOB mode
- Maximum LOB size (KB): 32
- Enable validation: (Description: Choose this setting if you want AWS DMS to compare the data at the source and the target, immediately after it performs a full data load. Validation ensures that your data was migrated accurately, but it requires additional time to complete.)
- Enable CloudWatch logs: (Description: CloudWatch logs usage will be charged at standard rates. See [here](#) for more details.)

Table mappings:

Editing mode:
 Guided UI: Set up your table mapping rules using a step-by-step guided interface.
 JSON editor: Learn more [\[\]](#). Enter your table mapping rules directly, in JSON format.

Specify at least one selection rule with an include action. After you do this, you can add one or more transformation rules.

Selection rules:

Choose the schema and/or tables you want to include with, or exclude from, your migration task. [Info](#) Add new selection rule

▼ where schema name is like 'DMS_SAMPLE' and table name is like '%', include X

AWS Database Migration Lab – Oracle Migration

Schema
DMS_SAMPLE

Table name
Use the % character as a wildcard
%

Action
Choose "Include" to migrate your selected objects, or "Exclude" to ignore them during the migration.
Include

Source filters [Info](#) Add column filter

▼ Transformation rules

You can use transformation rules to change or transform schema, table or column names of some or all of the selected objects. [Info](#) Add new transformation rule

▼ where schema name is like 'DMS_SAMPLE' and table name is like '%', convert-lowercase

Target Schema
Schema name DMS_SAMPLE

Action Make lowercase

▼ where schema name is like 'DMS_SAMPLE' and table name is like '%', convert-lowercase

Target Table
Schema name DMS_SAMPLE

Table name Use the % character as a wildcard
%

Action Make lowercase

▼ where schema name is like 'DMS_SAMPLE' and table name is like '%', convert-lowercase

Target Column
Schema name DMS_SAMPLE

Table name Use the % character as a wildcard
%

Column name %

Action Make lowercase

► Advanced task settings

[Cancel](#) [Create task](#)

33. After entering the values click on **Create task**. **** DMS task takes approximately 25 minutes to complete ****

34. At this point, the task should start migrating data from the source Oracle database to the Amazon Aurora RDS instance.

The screenshot shows the AWS DMS console with the 'Database migration tasks' page. The left sidebar has 'Conversion & migration' expanded, with 'Database migration tasks' selected. The main area displays a table with one row:

Name	Status	Source	Target	Type	Progress	Elapsed time
oracle-migration-task	Running	samplesource	auroratarget	Full load, ongoing replication	56 %	10 m

35. As the rows are being transferred, you can monitor the task progress:

- d. Click on your task (**oracle-migration-task**) and scroll to the **Table statistics** section to view the table statistics to see how many rows have been moved.
- e. If there is an error, the status color changes from green to red. Click on **View logs** link for the logs to debug.

Inspect the Content of Target Database

37. If you already disconnected from the EC2 Server, follow steps 1, 2, and 3 at the top of this lab to connect (RDP) to your EC2 instance.
38. Open **pgAdmin4** from within the EC2 server, and then connect to the Target Aurora RDS (PostgreSQL) database connection that you created earlier.
39. Inspect the migrated data, by querying one of the tables in the target database. For example, the following query should return a table with two rows:

```
SELECT *
FROM dms_sample.sport_type;
```

The screenshot shows the pgAdmin4 interface. In the left sidebar, under 'Servers (1)', there is one entry: 'Target Aurora RDS (PostgreSQL)'. Underneath it, 'Databases' contains 'AuroraDB'. The 'Schemas' section is expanded, showing five schemas: 'aws_oracle_data', 'aws_oracle_ext', 'aws_oracle_ext_keys', 'dms_sample', and 'public'. The 'Tables' section is collapsed. In the center, the 'Query Editor' tab is active, displaying the following SQL code:

```
1 SELECT *
2 FROM dms_sample.sport_type;
3
```

Below the query editor, the 'Data Output' tab is selected, showing the results of the query:

	name	description
1	baseball	A sport with 9 players, bats,...
2	football	Teams of 11 players attempt...

At the bottom right of the pgAdmin window, a green message box indicates: 'Successfully run. Total query runtime: 59 msec. 2 rows affected.'

Baseball, and football are the only two sports that are currently listed in this table.

Capture and Replicate Data Changes

40. Create another **Data Migration Task** with the following values for capturing data changes to the source Oracle database, and replicating the changes to the target Aurora RDS instance.

Parameter	Value
Task identifier	oracle-replication-task
Replication instance	oracle-replication
Source database endpoint	oracle-source
Target database endpoint	aurora-target
Migration type	Replicate data changes only
Start task on create	Checked
CDC stop mode	Don't use custom CDC stop mode
Target table preparation mode	Do nothing
Stop task after full load completes	Don't stop
Include LOB columns in replication	Limited LOB mode
Max LOB size (KB)	32
Enable validation	Unchecked
Enable CloudWatch logs	Checked

41. Expand the **Table mappings** section, and select **Guided UI** for the editing mode.

42. Add the same **Selection**, and **Transformation** rules as specified in steps 28, and 29.

AWS Database Migration Lab – Oracle Migration

The screenshot shows the AWS DMS console with the 'Create database migration task' wizard open. The left sidebar shows navigation options like Dashboard, Conversion & migration, and Database migration tasks. The main form is divided into sections: Task configuration, Task settings, Table mappings, and a summary step.

Task configuration:

- Task identifier: oracle-replication-task
- Replication instance: oracle-replication - vpc-0070bce6424a3e253
- Source database endpoint: oo1709diyhdpjb
- Target database endpoint: aurora-target
- Migration type: Replicate data changes only

A warning message in the configuration section states: "⚠ Your source database is Oracle. Replicating ongoing changes requires supplemental logging to be turned on. Please ensure your archive logs are retained on the server for a sufficient amount of time, (24 hours is usually enough.) To set your archivelog retention on RDS databases you can use the following command: exec rdsadmin:rdsadmin_util.set_configuration('archivelog retention hours', 24);".

Start task on create

Task settings:

- CDC start mode: Don't use custom CDC start mode
- CDC stop mode: Don't use custom CDC stop mode
- Target table preparation mode: Do nothing
- Include LOB columns in replication: Limited LOB mode
- Maximum LOB size (KB): 32
- Enable validation: Choose this setting if you want AWS DMS to compare the data at the source and the target, immediately after it performs a full data load. Validation ensures that your data was migrated accurately, but it requires additional time to complete.
- Enable CloudWatch logs: CloudWatch logs usage will be charged at standard rates. See [here](#) for more details.

Table mappings:

Editing mode:
 Guided UI: Set up your table mapping rules using a step-by-step guided interface.
 JSON editor: Learn more [\[\]](#). Enter your table mapping rules directly, in JSON format.

Specify at least one selection rule with an include action. After you do this, you can add one or more transformation rules.

AWS Database Migration Lab – Oracle Migration

▼ Selection rules

Choose the schema and/or tables you want to include with, or exclude from, your migration task. [Info](#) [Add new selection rule](#)

▼ where schema name is like 'DMS_SAMPLE' and table name is like '%', include [X](#)

Schema
DMS_SAMPLE

Table name
Use the % character as a wildcard
%

Action
Choose "Include" to migrate your selected objects, or "Exclude" to ignore them during the migration.
Include

Source filters [Info](#) [Add column filter](#)

▼ Transformation rules

You can use transformation rules to change or transform schema, table or column names of some or all of the selected objects. [Info](#) [Add new transformation rule](#)

▼ where schema name is like 'DMS_SAMPLE' and table name is like '%', convert-lowercase [X](#)

Target
Schema

Schema name
DMS_SAMPLE

Action
Make lowercase

▼ where schema name is like 'DMS_SAMPLE' and table name is like '%', convert-lowercase [X](#)

Target
Table

Schema name
DMS_SAMPLE

Table name
Use the % character as a wildcard
%

Action
Make lowercase

▼ where schema name is like 'DMS_SAMPLE' and table name is like '%', convert-lowercase [X](#)

Target
Column

Schema name
DMS_SAMPLE

Table name
Use the % character as a wildcard
%

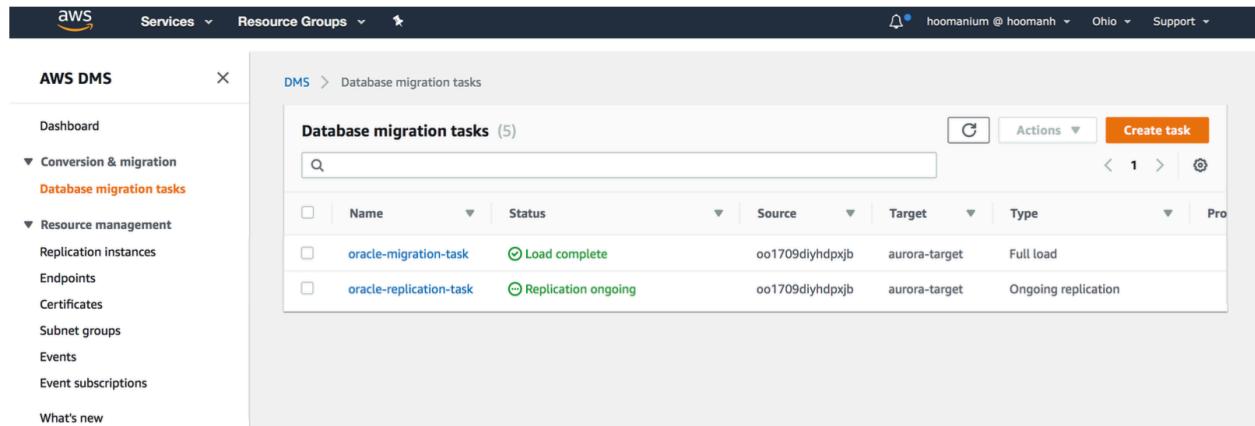
Column name
%

Action
Make lowercase

► Advanced task settings

[Cancel](#) [Create task](#)

43. After entering the values click on **Create task**. **** Task takes approximately 2 minutes to run ****
44. At this point, the new migration task is ready to replicate ongoing data changes from the source Oracle RDS to the Amazon Aurora RDS (PostgreSQL) database.



The screenshot shows the AWS DMS console with the 'Database migration tasks' page. The left sidebar has 'Conversion & migration' expanded, with 'Database migration tasks' selected. The main area displays a table of migration tasks:

Name	Status	Source	Target	Type
oracle-migration-task	Load complete	oo1709diyhdpzb	aurora-target	Full load
oracle-replication-task	Replication ongoing	oo1709diyhdpzb	aurora-target	Ongoing replication

Replicating Changes from Source to the Target Database

Now you are going to simulate a transaction to the source database by updating the **sport_type** table. The Database Migration Service will automatically detect and replicate these changes to the target database.

45. Use **Oracle SQL Developer** connect to the source Oracle RDS instance (described in steps 4, and 5.)

46. Open a **New Query** window and **execute** the following statement to insert 5 new sports into the **sport_type** table:

```
INSERT ALL
```

```
INTO dms_sample.sport_type (name,description) VALUES ('hockey', 'A sport in which two teams play against each other by trying to move a puck into the opponents goal using a hockey stick')
```

```
INTO dms_sample.sport_type (name,description) VALUES ('basketball', 'A sport in which two teams of five players each that oppose one another shoot a basketball through the defenders hoop')
```

```
INTO dms_sample.sport_type (name,description) VALUES ('soccer','A sport played with a spherical ball between two teams of eleven players')
```

```
INTO dms_sample.sport_type (name,description) VALUES ('volleyball','two teams of six players are separated by a net and each team tries to score by grounding a ball on the others court')
```

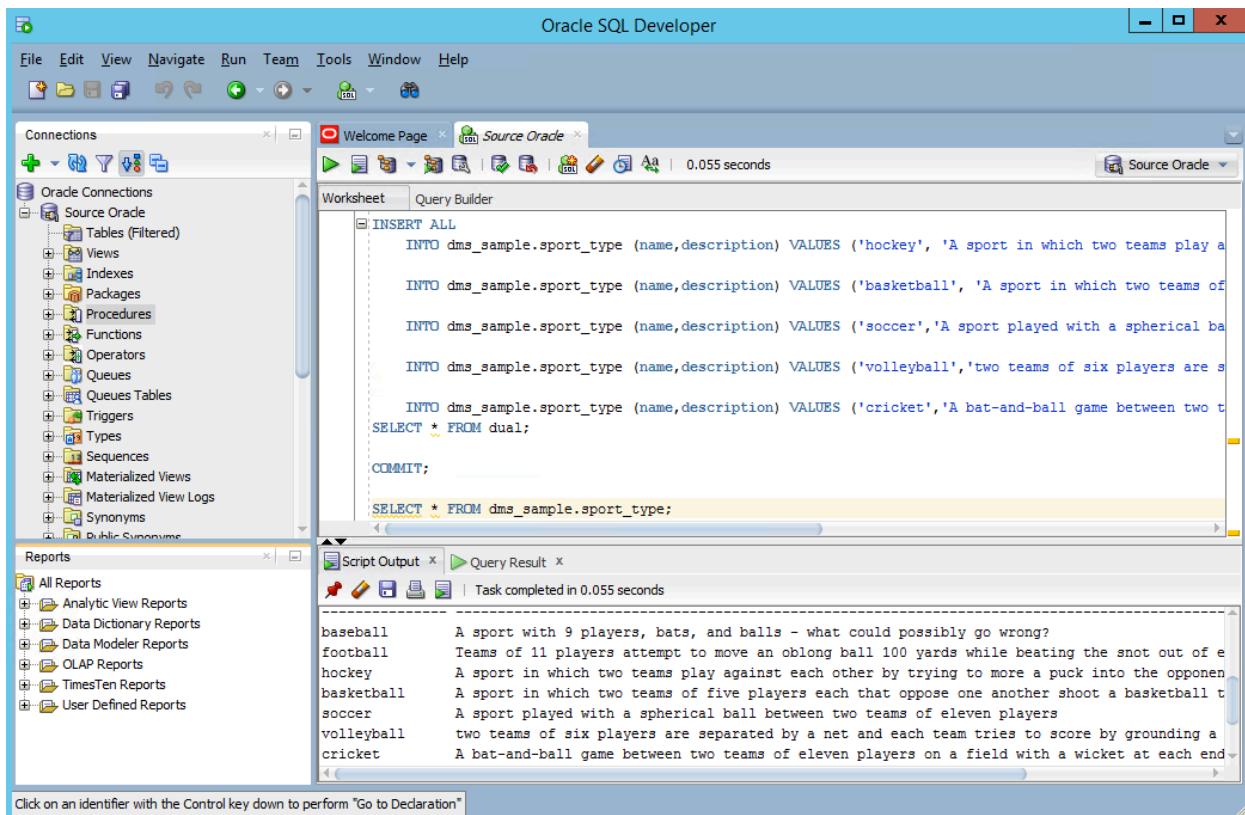
```
INTO dms_sample.sport_type (name,description) VALUES ('cricket','A bat-and-ball game between two teams of eleven players on a field with a wicket at each end')
```

```
SELECT * FROM dual;
```

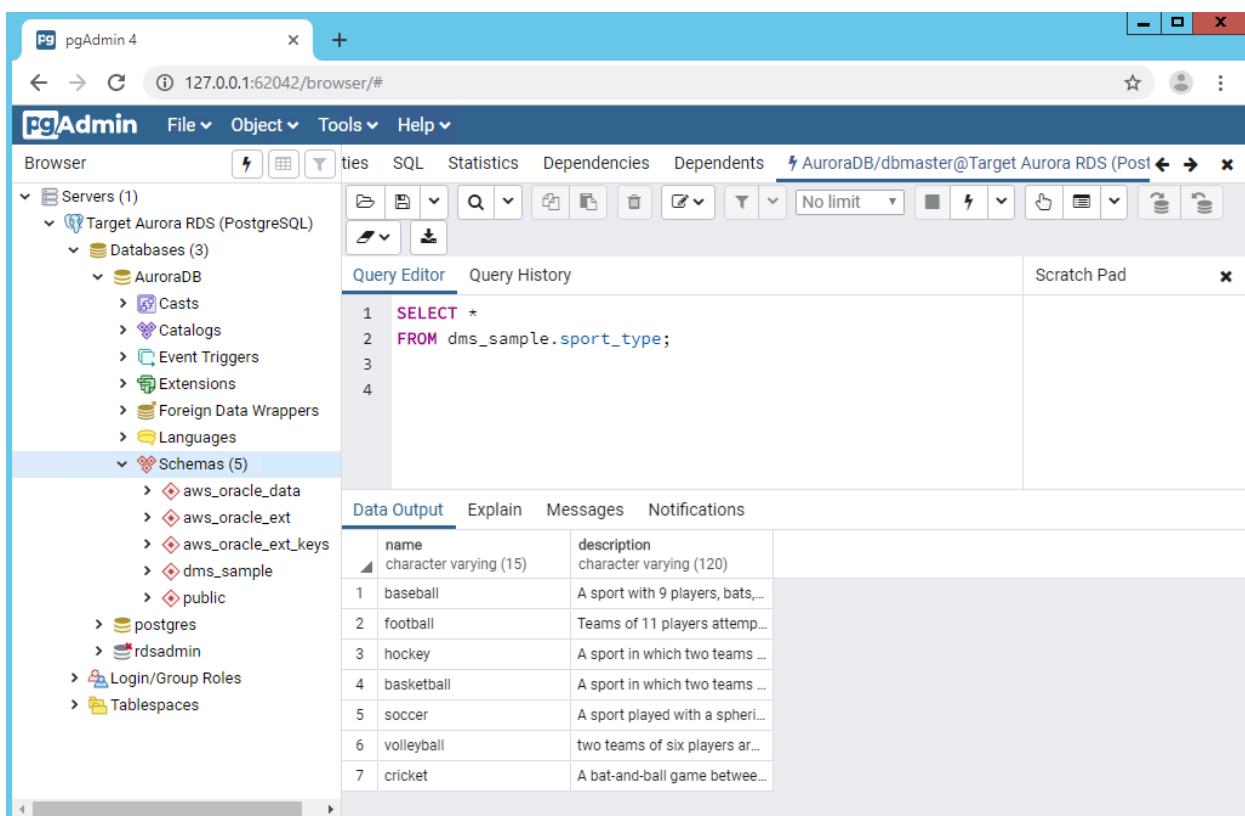
```
COMMIT;
```

```
SELECT * FROM dms_sample.sport_type;
```

AWS Database Migration Lab – Oracle Migration



47. Repeat steps 35 and 36 to inspect the content of **sport_type** table in the target database.



Notice the new records for: basketball, cricket, hokey, soccer, and volleyball that you added to the **sports_type** table in the source database have been replicated to your **dms_reporting** database. You can further investigate the number of inserts, deletes, updates, and DDLs by viewing the **Table statistics** of your **Database migration tasks** in AWS console.

The AWS DMS task keeps the target Aurora PostgreSQL database up to date with source database changes. AWS DMS keeps all the tables in the task up to date until it's time to implement the application migration. The latency is close to zero when the target has caught up to the source.

Summary

The first part of the lab demonstrated how easy it is to migrate the data from an Oracle database running on an Amazon RDS instance to an Amazon Aurora (PostgreSQL) instance using the AWS Database Migration Service (DMS). Similarly, you observed how DMS automatically replicates new transactions on the source to target database.

You can follow the same steps to migrate SQL Server and Oracle workloads to other RDS engines including PostgreSQL and MySQL.